



UNIVERSIDADE*FEDERAL
DE*MINAS*GERAIS

Information Retrieval

Introduction

Rodrygo L. T. Santos
rodrygo@dcc.ufmg.br

Information retrieval

“

Information retrieval is a field concerned with the structure, analysis, organization, storage, searching, and retrieval of information.

◦ Gerard Salton, 1968

Retrieval tasks





Google

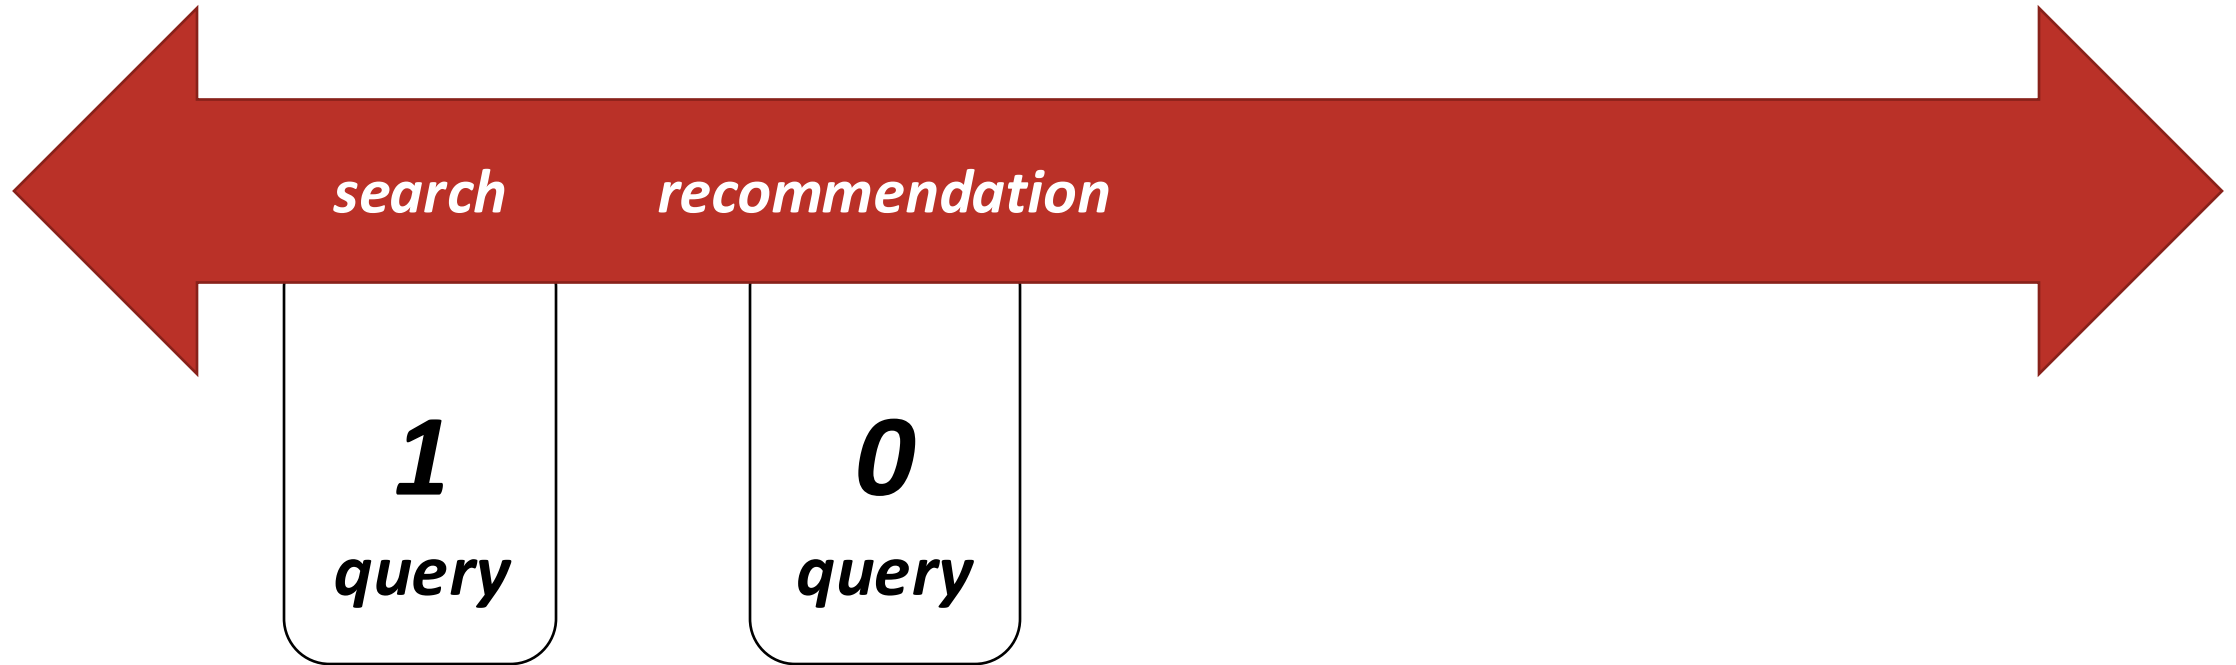


Google Search

I'm Feeling Lucky

Google offered in: [Português \(Brasil\)](#)

Retrieval tasks





Hi, Rodrygo

CUSTOMER SINCE 2011

YOUR ORDERS

0 recent orders

TOP CATEGORIES FOR YOU

Musical Instruments

Electronics

Toys & Games

PRIME

FAST, FREE SHIPPING

50 million eligible items



FRESH

SHOP GROCERY DEALS

30-day FREE trial



VIDEO

INCLUDED WITH PRIME

Top movies & TV shows

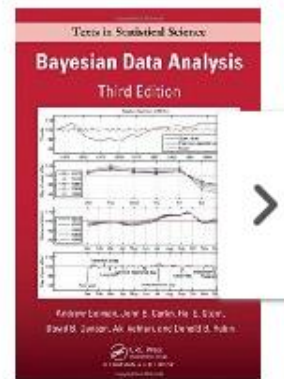
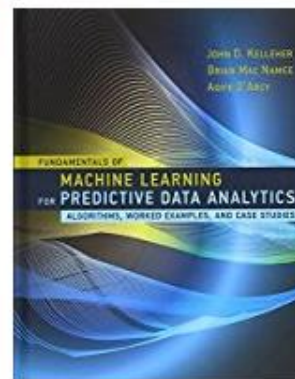
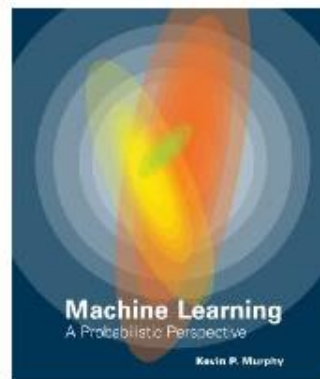
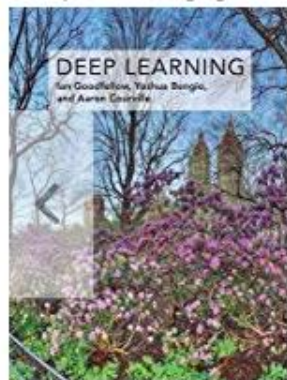


M

AMA



Inspired by your shopping trends

More convenience
Prime Delivery

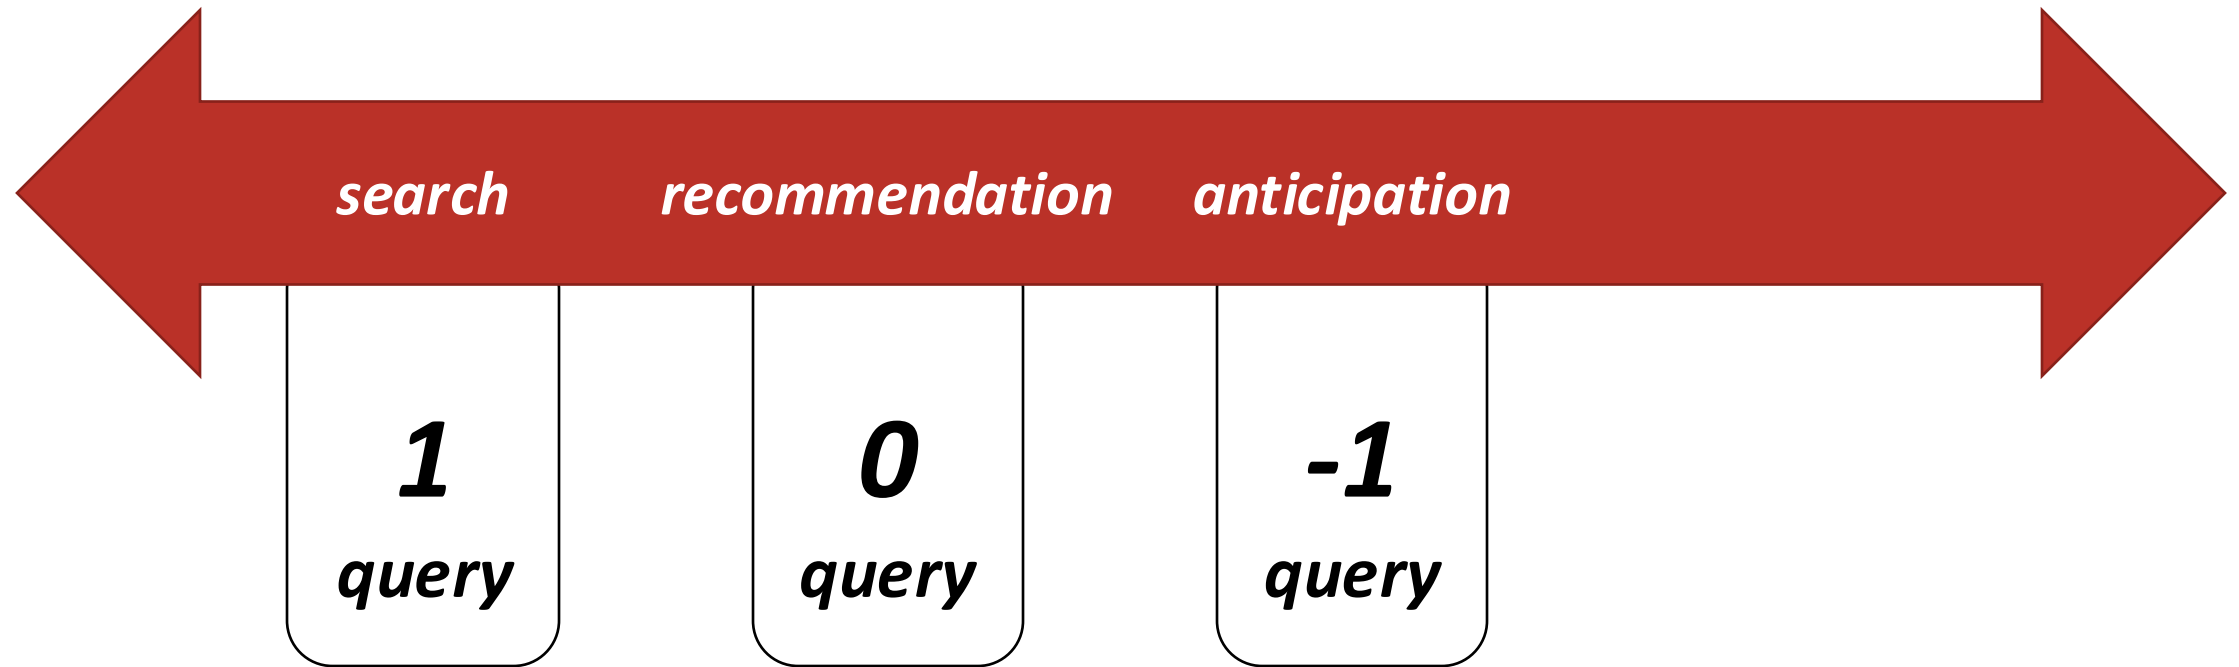
Recommendations for you in Books



Ad feedback



Retrieval tasks




Gate
51a

Terminal
2

Category

Ticket type: World Perk Rewards Premier Access

 [View email](#)

Car rental

Economy 2 door sedan

Hertz rental car reservation

Name
Mr. John Smith

Booking Number
E12345678

Thu, 18 Apr, 2013
11:40

Fri, 26 Apr, 2013
21:50

Hertz San Diego
987 Harbor Dr, San Diego, CA 92101

 [Get directions](#)

 [Manage reservation](#)

 [View email](#)

Next Appointment

Agency Meeting

11:30 AM

Ninth Ave, New York, NY 10011

 [Email guests](#)

Flights

Delta Air Lines

San Francisco



63°

SCATTERED
CLOUDS
5mph
10%

TUE
68°
48°

WED
67°
44°

THU
65°
48°

FRI
57°
46°

Hotels

The Connaught Hotel

Carlos Place, Mayfair, London W1K 2AL, United Kingdom

Check in from 12:00pm today



 [Call](#)

 [Hotel information](#)

 [View email](#)

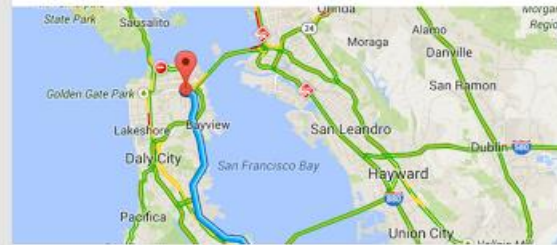
Packages



Traffic & Transit

57 mins to work

Normal traffic on US - 101



 [Navigate / 57 mins via US - 101](#)

Restaurant Reservations

Broder

2508 SE Clinton St, Portland, OR 97202

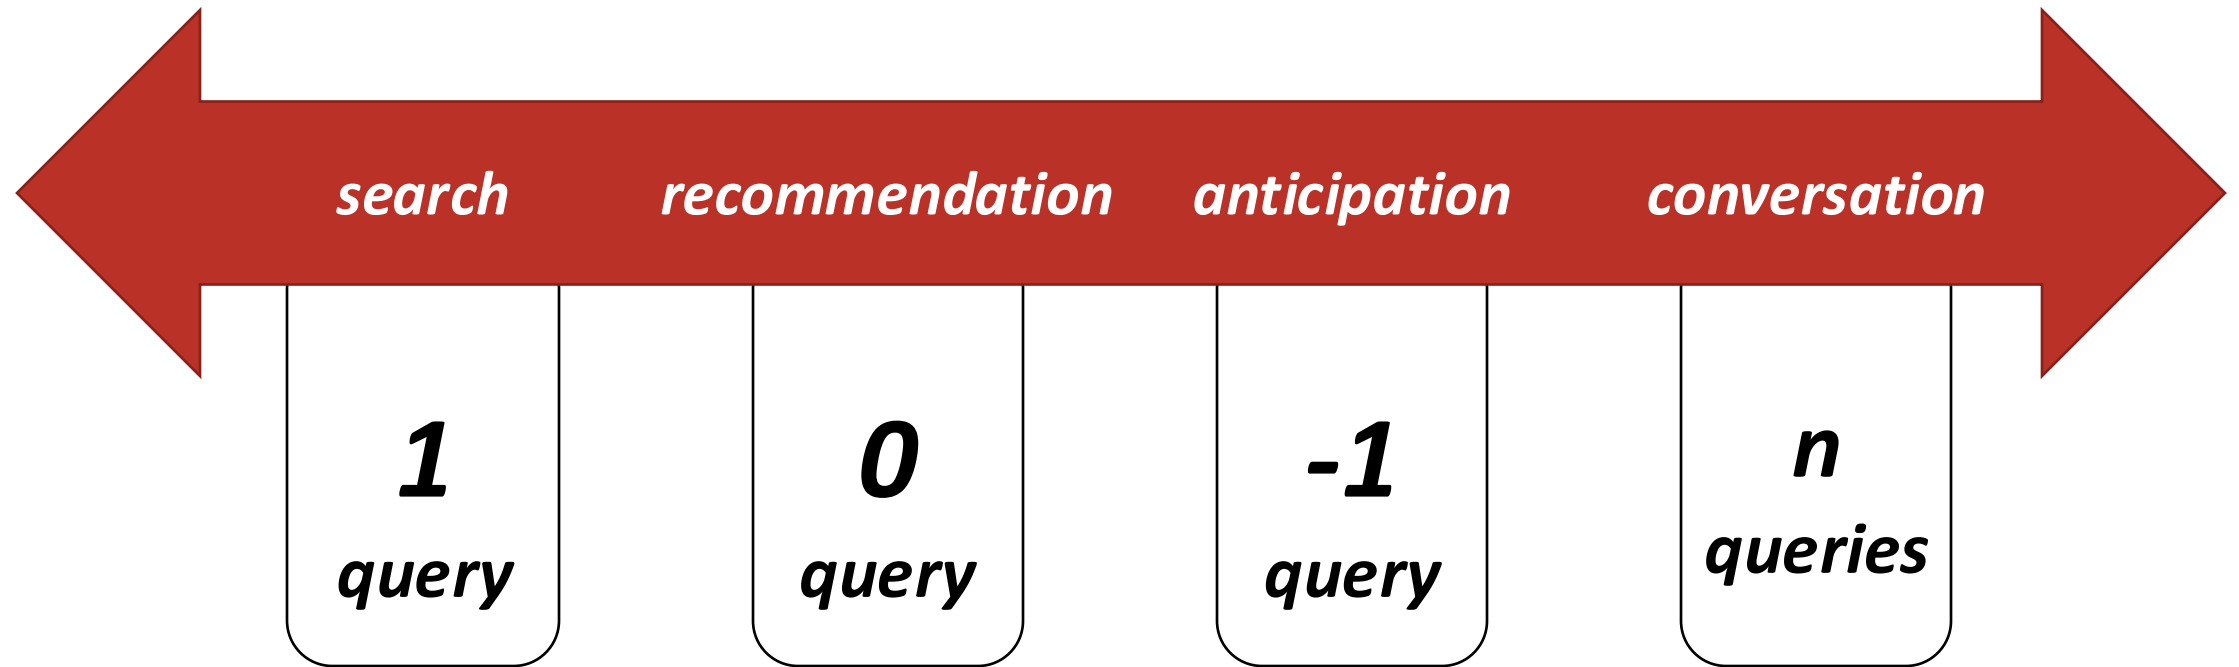
Reservation in 1 hour

Travel time walking 45 minutes



 [Get directions](#)

Retrieval tasks





ChatGPT 

Memory Full 

 Temporary



ChatGPT



Explore GPTs

What can I help with?

what is information retrieval|



Search



Reason



what is information retrieval and its applications?

what is information retrieval and how does it work in search engines?

what is information retrieval and how does it work

what is information retrieval and its role in search engines?



Upgrade plan

More access to the best models

ChatGPT can make mistakes. Check important info.

?



Today

Information Retrieval Overview

what is information retrieval

Information Retrieval (IR) is the process of obtaining relevant information from a large collection of data, typically in response to a user's query. The goal is to efficiently retrieve documents, web pages, images, or other types of information that match the user's needs. IR systems are widely used in search engines (e.g., Google, Bing), document retrieval (e.g., academic search engines like Google Scholar), and recommender systems.

Key Aspects of Information Retrieval:

1. **Indexing** – Preprocessing data to create a structured representation for fast searching.
2. **Query Processing** – Interpreting and optimizing user queries to improve retrieval performance.
3. **Ranking** – Scoring and ordering retrieved items based on relevance, often using machine-learned ranking models.
4. **Evaluation** – Measuring retrieval effectiveness using metrics like precision, recall, and NDCG (Normalized Discounted Cumulative Gain).

Your research in recommender systems overlaps significantly with IR, especially in ranking algorithms and leveraging contextual information for better retrieval. Let me know if you want me to focus on any particular IR concept!

explain it to a toddler

Imagine you have a big toy box filled with all kinds of toys. Now, you want to find your favorite red car.

2. **Query Processing** – Interpreting and optimizing user queries to improve retrieval performance.
3. **Ranking** – Scoring and ordering retrieved items based on relevance, often using machine-learned ranking models.
4. **Evaluation** – Measuring retrieval effectiveness using metrics like precision, recall, and NDCG (Normalized Discounted Cumulative Gain).

Your research in recommender systems overlaps significantly with IR, especially in ranking algorithms and leveraging contextual information for better retrieval. Let me know if you want me to focus on any particular IR concept!

explain it to a toddler

Imagine you have a big toy box filled with all kinds of toys. Now, you want to find your favorite red car. Instead of digging through everything, you ask a magic robot helper:

"Hey, find my red car!"

The robot quickly looks through all the toys, finds the red car, and gives it to you.

Information Retrieval is like that magic robot, but instead of toys, it finds information—like pictures, songs, or answers to questions—when you search for them on the internet! 🚀



Ask anything



Search



Reason



Upgrade plan

More access to the best models

Our focus: search

User's perspective: accessing information

- Relevance, speed

Advertiser's perspective: publicity

- More attention, less ad costs

Search engine's perspective: monetization

- More users, more ad revenue, less op costs

**What does a
search
engine do?**

Spelling correction



- information **retrieval**
- information **retrieval is fastest from**
- information **retrieval system pdf**
- information **retrieval services**
- information **retrieval in nlp**
- information **retrieval models**
- information **retrieval meaning**
- information **retrieval course**
- information **retrieval system notes pdf**
- information **retrieval book**

[Report inappropriate predictions](#)

Query autocompletion
Query suggestions

All Images Videos Short videos News Shopping Forums : More

Refine results



Under R\$2,000

Vertical search results

Sort by

Price: low to high

Price: high to low

Stores

☐ eBay

☐ Mercado Livre

☐ Ubuy

☐ Amazon.com.br

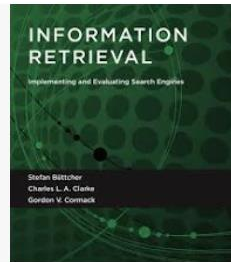
[See 20 more](#)

Price

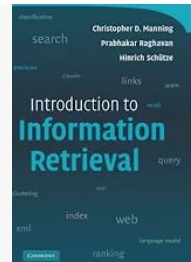
Under R\$2,000

Over R\$2,000

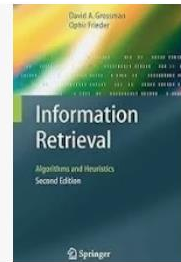
R\$ Min R\$ M... Go



Information Retrieval: Implementing and...
R\$254.26 now R\$50....
 Amazon.c... & more



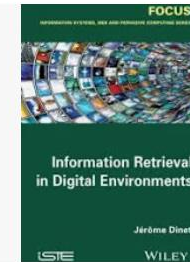
Introduction to Information Retrieval
R\$389.52 (£52)
 eBooks.com & more
5.0 ★★★★★ (1)



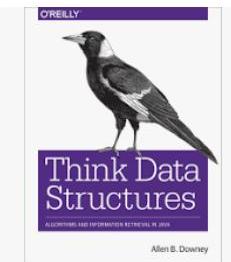
Information Retrieval: Algorith...
R\$470.00 Pre-owned
 Nanah Cul... & more
1.0 ★☆☆☆☆ (1)



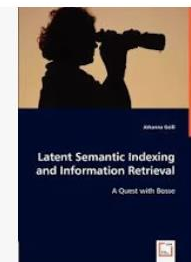
Information Retrieval: Data Structures &...
R\$42.70 Pre-owned
 eBay - thrift.books



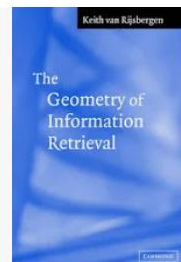
Information Retrieval in Digital Environments
R\$1,025.48 (\$178)
 Wiley
30-day returns



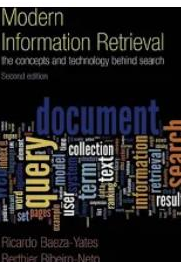
Think Data Structures: Algorithms and...
R\$196.36 now R\$65....



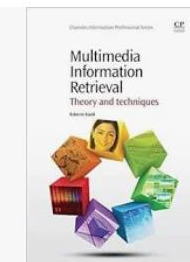
Latent Semantic Indexing And Information Retrieval
R\$500.15 now R\$50....



The Geometry of Information Retrieval
R\$757.23 now R\$25...
 Umlivro



Modern Information Retrieval: The...
R\$813.07 (\$141)
 eBay - gra... & more



Multimedia Information Retrieval: Theory...
R\$636.15 (SAR 414)



Search Labs | AI Overview

Learn more ⋮

Information Retrieval (IR) is the process of finding and accessing relevant information from a collection of information, often text-based, using search queries or other methods. It involves organizing, storing, and evaluating information to facilitate efficient access and retrieval. [↗](#)

Here's a more detailed explanation:

What is Information Retrieval? [↗](#)

- IR is a field concerned with finding, organizing, and retrieving information from large collections, especially in a structured format like text. [↗](#)
- It involves tasks such as indexing, querying, and presenting information in a way

Show more ▼

What Is Information Retrieval?

May 15, 2024



Coveo



What is Information Retrieval? - GeeksforGeeks

Sep 19, 2023 — What is Information Retrieval? * Information Retrieval (IR) can be defined as a software program that deal...



GeeksforGeeks



Generated answers



Wikipedia

https://en.wikipedia.org/wiki/Information_retrieval

Information retrieval

Information retrieval (IR) in computing and information science is the task of **identifying and retrieving information system resources** that are relevant to ...

[Music information retrieval](#)

[Information needs](#)

[Boolean model of information...](#)

People also ask :

What do you mean by information retrieval?



What is an example of retrieval information?



What is information retrieval in NLP?



What are the three types of information retrieval?



[Feedback](#)



Stanford University

<https://nlp.stanford.edu/IR-book/information-retrieva...>

Introduction to Information Retrieval

The book aims to **provide a modern approach to information retrieval** from a computer science perspective. It is based on a course we have been teaching in ...



GeeksforGeeks

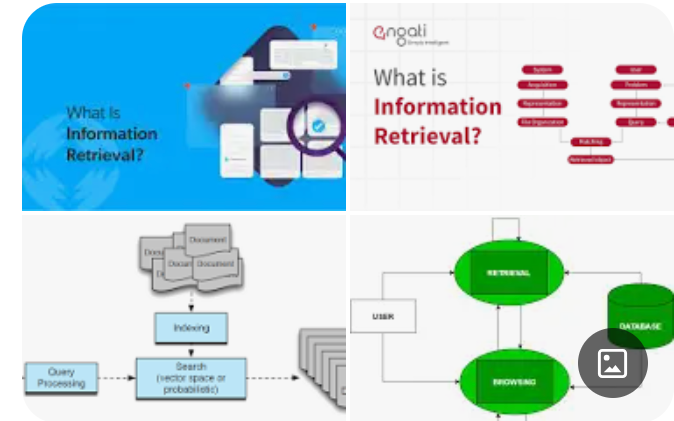
<https://www.geeksforgeeks.org/what-is-information-re...>

What is Information Retrieval?

Sep 19, 2023 — **Information Retrieval** is the activity of obtaining material that can usually be documented on an unstructured nature i.e. usually text which ...



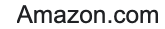
Information retrieval



Information retrieval in computing and information science is the task of identifying and retrieving information system resources that are relevant to an information need. The information need can be specified in the form of a search query.

Source: [Wikipedia](#)

Knowledge snippets



⋮

search

Christopher D. Manning
Instructor, Stanford University

Introduction to
**Information
Retrieval**

search engine

document

query

index

retrieval

algorithm

model

data

system

language

text

web

information

computer

science

technology

research

education

learning

artificial intelligence

natural language processing

machine learning

deep learning

neural networks

convolutional neural networks

recurrent neural networks

word embeddings

topic modeling

latent Dirichlet allocation

probabilistic graphical models

Bayesian networks

Markov chains

Monte Carlo methods

stochastic gradient descent

backpropagation

softmax

cross entropy

F1 score

precision

recall

area under the curve

ROC curve

confusion matrix

accuracy

specificity

sensitivity

positive predictive value

negative predictive value

positive likelihood ratio

negative likelihood ratio

diagnostic odds ratio

number needed to treat

number needed to harm

relative risk

odds ratio

hazard ratio

incidence rate

prevalence

incidence

mortality

morbidity

quality of life

health economics

health services research

public health

epidemiology

biostatistics

genetics

genomics

proteomics

metabolomics

transcriptomics

epigenomics

microbiomics

single-cell genomics

CRISPR

gene editing

synthetic biology

nanotechnology

robotics

artificial intelligence

machine learning

deep learning

neural networks

convolutional neural networks

recurrent neural networks

word embeddings

topic modeling

latent Dirichlet allocation

probabilistic graphical models

Bayesian networks

Markov chains

Monte Carlo methods

stochastic gradient descent

backpropagation

softmax

cross entropy

F1 score

precision

recall

area under the curve

ROC curve

confusion matrix

accuracy

specificity

sensitivity

positive predictive value

negative predictive value

positive likelihood ratio

negative likelihood ratio

diagnostic odds ratio

number needed to treat

number needed to harm

relative risk

odds ratio

hazard ratio

incidence rate

prevalence

incidence

mortality

morbidity

quality of life

health economics

health services research

public health

epidemiology

biostatistics

genetics

genomics

proteomics

metabolomics

transcriptomics

epigenomics

microbiomics

single-cell genomics

CRISPR

gene editing

synthetic biology

nanotechnology

robotics

artificial intelligence

machine learning

deep learning

neural networks

convolutional neural networks

recurrent neural networks

word embeddings

topic modeling

latent Dirichlet allocation

probabilistic graphical models

Bayesian networks

Markov chains

Monte Carlo methods

stochastic gradient descent

backpropagation

softmax

cross entropy

F1 score

precision

recall

area under the curve

ROC curve

confusion matrix

accuracy

specificity

sensitivity

positive predictive value

negative predictive value

positive likelihood ratio

negative likelihood ratio

diagnostic odds ratio

number needed to treat

number needed to harm

relative risk

odds ratio

hazard ratio

incidence rate

prevalence

incidence

mortality

morbidity

quality of life

health economics

health services research

public health

epidemiology

biostatistics

genetics

genomics

proteomics

metabolomics

transcriptomics

epigenomics

microbiomics

single-cell genomics

CRISPR

gene editing

synthetic biology

nanotechnology

robotics

artificial intelligence

machine learning

deep learning

neural networks

convolutional neural networks

recurrent neural networks

word embeddings

topic modeling

latent Dirichlet allocation

probabilistic graphical models

Bayesian networks

Markov chains

Monte Carlo methods

stochastic gradient descent

backpropagation

softmax

cross entropy

F1 score

precision

recall

area under the curve

ROC curve

confusion matrix

accuracy

specificity

sensitivity

positive predictive value

negative predictive value

positive likelihood ratio

negative likelihood ratio

diagnostic odds ratio

number needed to treat

number needed to harm

relative risk

odds ratio

hazard ratio

incidence rate

prevalence

incidence

mortality

morbidity

quality of life

health economics

health services research

public health

epidemiology

biostatistics

genetics

genomics

proteomics

metabolomics

transcriptomics

epigenomics

microbiomics

single-cell genomics

CRISPR

gene editing

synthetic biology

nanotechnology

robotics

artificial intelligence

machine learning

deep learning

neural networks

convolutional neural networks

recurrent neural networks

word embeddings

topic modeling

latent Dirichlet allocation

probabilistic graphical models

Bayesian networks

Markov chains

Monte Carlo methods

stochastic gradient descent

backpropagation

softmax

cross entropy

F1 score

precision

recall

area under the curve

ROC curve

confusion matrix

accuracy

specificity

sensitivity

positive predictive value

negative predictive value

positive likelihood ratio

negative likelihood ratio

diagnostic odds ratio

number needed to treat

number needed to harm

relative risk

odds ratio

hazard ratio

incidence rate

prevalence

incidence

mortality

morbidity

quality of life

health economics

health services research

public health

epidemiology

biostatistics

genetics

genomics

proteomics

metabolomics

transcriptomics

epigenomics

microbiomics

single-cell genomics

CRISPR

gene editing

synthetic biology

nanotechnology

robotics

artificial intelligence

machine learning

deep learning

neural networks

convolutional neural networks

recurrent neural networks

word embeddings

topic modeling

latent Dirichlet allocation

probabilistic graphical models

Bayesian networks

Markov chains

Monte Carlo methods

stochastic gradient descent

backpropagation

softmax

cross entropy

F1 score

precision

recall

area under the curve

ROC curve

confusion matrix

accuracy

specificity

sensitivity

positive predictive value

negative predictive value

positive likelihood ratio

negative likelihood ratio

diagnostic odds ratio

number needed to treat

number needed to harm

relative risk

odds ratio

hazard ratio

incidence rate

prevalence

incidence

mortality

morbidity

quality of life

health economics

health services research

public health

epidemiology

biostatistics

genetics

genomics

proteomics

metabolomics

transcriptomics

epigenomics

microbiomics

single-cell genomics

CRISPR

gene editing

synthetic biology

nanotechnology

robotics

artificial intelligence

machine learning

deep learning

neural networks

convolutional neural networks

recurrent neural networks

word embeddings

topic modeling

latent Dirichlet allocation

probabilistic graphical models

Bayesian networks

Markov chains

Monte Carlo methods

stochastic gradient descent

backpropagation

softmax

cross entropy

F1 score

precision

recall

area under the curve

ROC curve

confusion matrix

accuracy

specificity

sensitivity

positive predictive value

negative predictive value

positive likelihood ratio

negative likelihood ratio

diagnostic odds ratio

number needed to treat

number needed to harm

relative risk

odds ratio

hazard ratio

incidence rate

prevalence

incidence

mortality

morbidity

quality of life

health economics

health services research

public health

epidemiology

biostatistics

genetics

genomics

proteomics

metabolomics

transcriptomics

epigenomics

microbiomics

single-cell genomics

CRISPR

gene editing

synthetic biology

nanotechnology

robotics

artificial intelligence

machine learning

deep learning

neural networks

convolutional neural networks

recurrent neural networks

word embeddings

topic modeling

latent Dirichlet allocation

probabilistic graphical models

Bayesian networks

Markov chains

Monte Carlo methods

stochastic gradient descent

backpropagation

softmax

cross entropy

F1 score

precision



⋮

n, 10 Mar 2025



Next

Class-tested and coherent, this groundbreaking new textbook teaches web-era information retrieval, including web search and the related areas of text ...
\$59.99 · 4.5 ★★★★★ (199)



arXiv

<https://arxiv.org> › [list](#) › [cs.IR](#) › recent ⋮

Information Retrieval

Information Retrieval. Authors and titles for recent submissions. Tue, 11 Mar 2025 · Mon, 10 Mar 2025 · Fri, 7 Mar 2025 · Thu, 6 Mar 2025 · Wed, 5 Mar 2025.

People also search for ⋮

Information retrieval **example**



Information retrieval **book**



Information retrieval **system**



Information retrieval **process**



Information retrieval **ppt**



Information retrieval **methods**



Information retrieval **pdf**



Information retrieval **model**



Results are personalised - [Try without personalisation](#)

Personalization

Brazil

● 31275 - Pampulha, Belo Horizonte - State of Minas Gerais - [Based on your places \(Work\)](#) - [Update location](#)

[Help](#)

[Send feedback](#)

[Privacy](#)

[Terms](#)



Wikipedia

https://en.wikipedia.org/wiki/Information_retrieval ⋮

Information retrieval

Information retrieval (IR) in computing and information science is the task of **identifying and retrieving information system resources** that are relevant to ...

[Music information retrieval](#)

[Information needs](#)

[Boolean model of information...](#)

People also ask ⋮

What do you mean by information retrieval?



What is an example of retrieval information?



What is information retrieval in NLP?



What are the three types of information retrieval?



[Feedback](#)



Stanford University

<https://nlp.stanford.edu/IR-book/information-retrieva...> ⋮

Introduction to Information Retrieval

The book aims to **provide a modern approach to information retrieval** from a computer science perspective. It is based on a course we have been teaching in ...



GeeksforGeeks

<https://www.geeksforgeeks.org/what-is-information-re...> ⋮

What is Information Retrieval?

Sep 19, 2023 — **Information Retrieval** is the activity of obtaining material that can usually be documented on an unstructured nature i.e. usually text which ...



ten
blue
links

Videos ⋮

The search problem

Given

- Some evidence of the user's need

Produce

- Relevant information

The search problem

Given

- Some evidence of the user's need

Produce

- A list of matching information items
- In decreasing order of relevance

The search problem

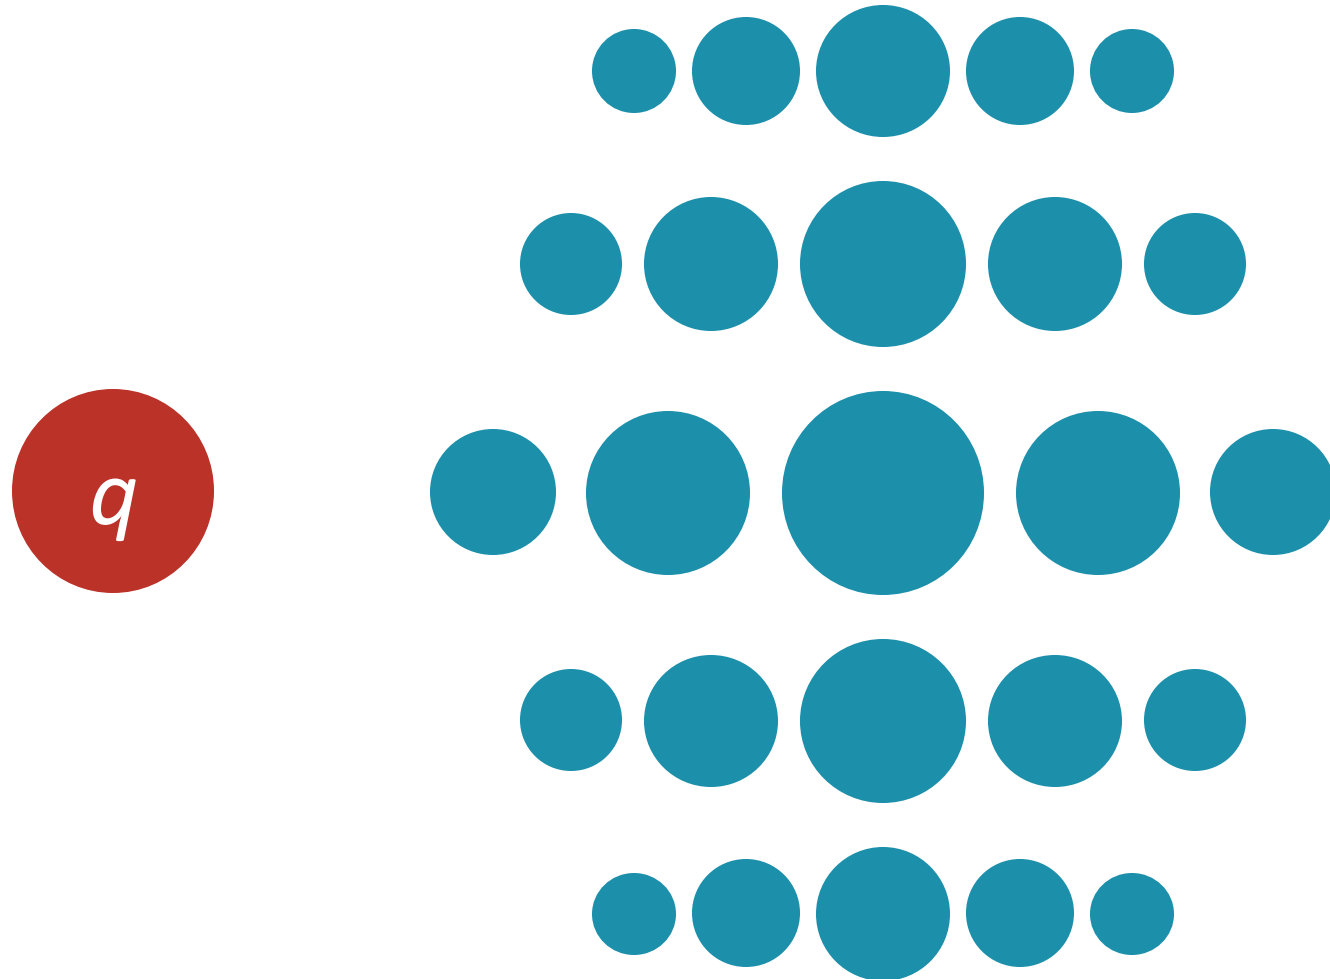
Given

- Some ~~evidence of the user's need~~ *query*

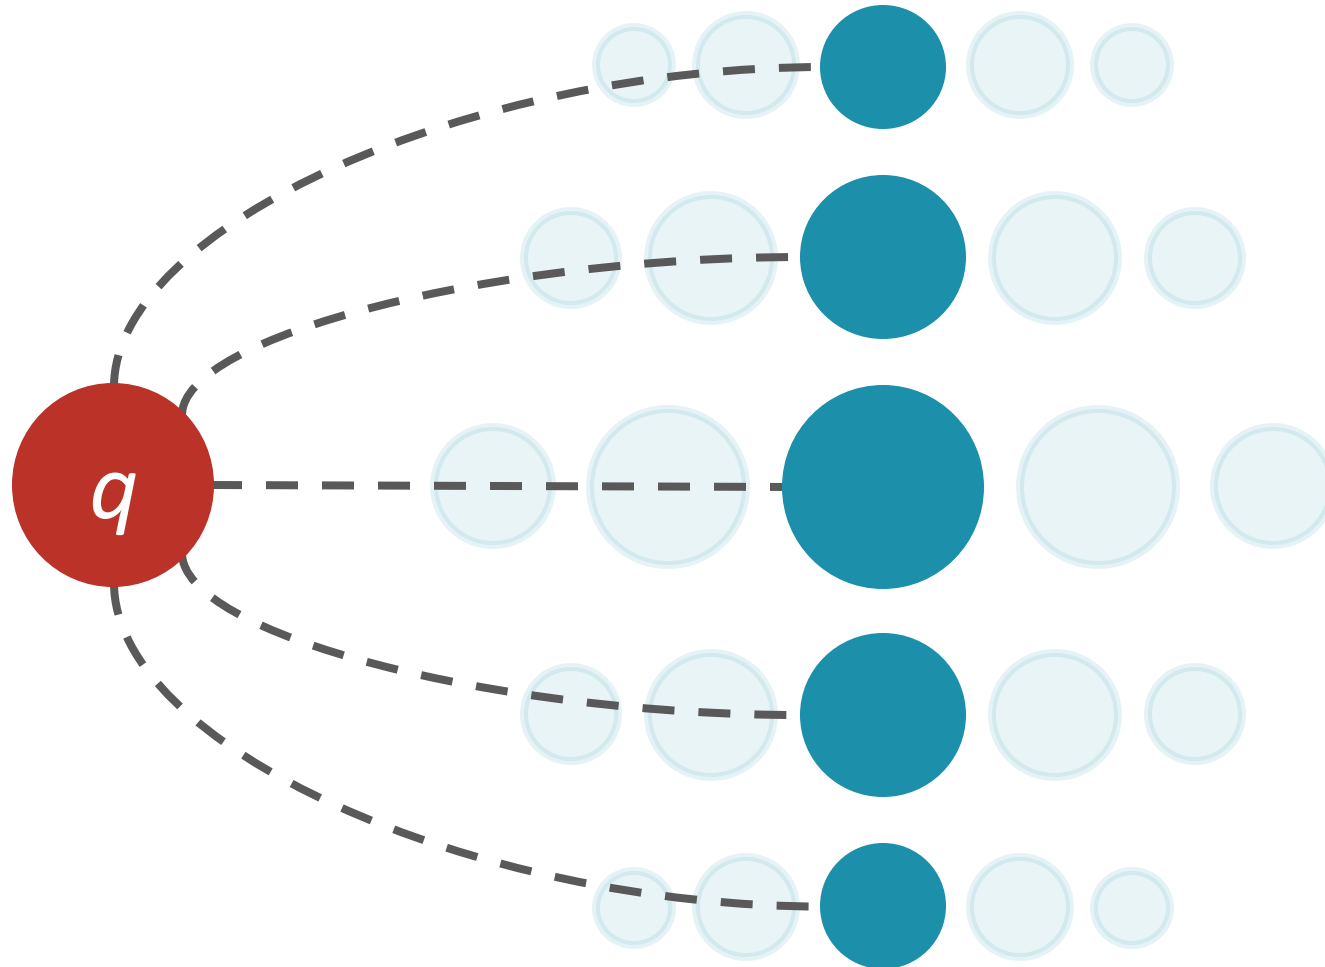
Produce

- A list of matching ~~information items~~ *documents*
- In decreasing order of relevance

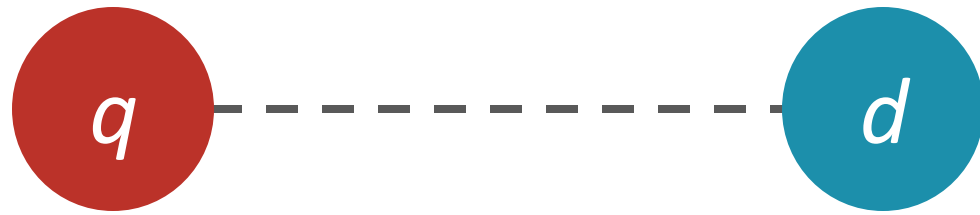
1) What documents do we show?



2) What order do we show them in?



2) What order do we show them in?



$f(q, d)$

**Isn't it a
solved
problem?**

A large, illuminated sign is mounted on a brick wall at night. The sign is white with black text. The word 'ChatGPT' is enclosed in a light pink, hand-drawn rectangular box. The rest of the text, 'DOESN'T HAVE ALL THE ANSWERS', is in a bold, sans-serif font. In the foreground, there are green plants. The background is dark, suggesting it is nighttime.

ChatGPT

DOESN'T
HAVE ALL THE
ANSWERS

Search in numbers

A lot of people *10^4 queries per second*

From a lot of places *whole planet (and beyond?)*

Using a lot of devices *smart-you-name-it*

Looking for a lot of info *10^{11} documents*

Spread all over the Internet *10^7 servers*

Efficiency

Efficiency is about doing something (good or bad) in an optimal way (i.e., faster or with fewer resources)

Key performance indicators

- *Query latency*: searching billions of documents
- *Query throughput*: serving thousands of users
- *Document latency*: serving freshly published content

Effectiveness

Effectiveness is about doing the right thing; it's about finding documents that are relevant to the user

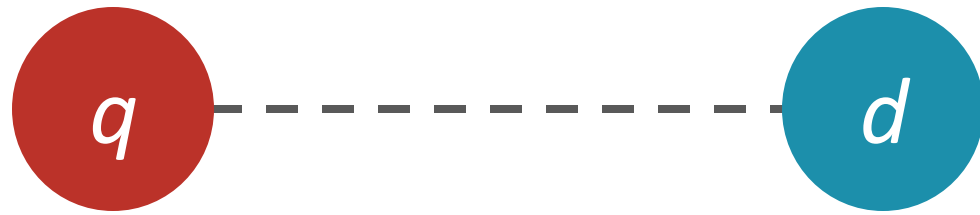
Relevance is influenced by many factors

- Topical relevance vs. user relevance
- Task, context, novelty, style

Ranking models define ***a view of*** relevance

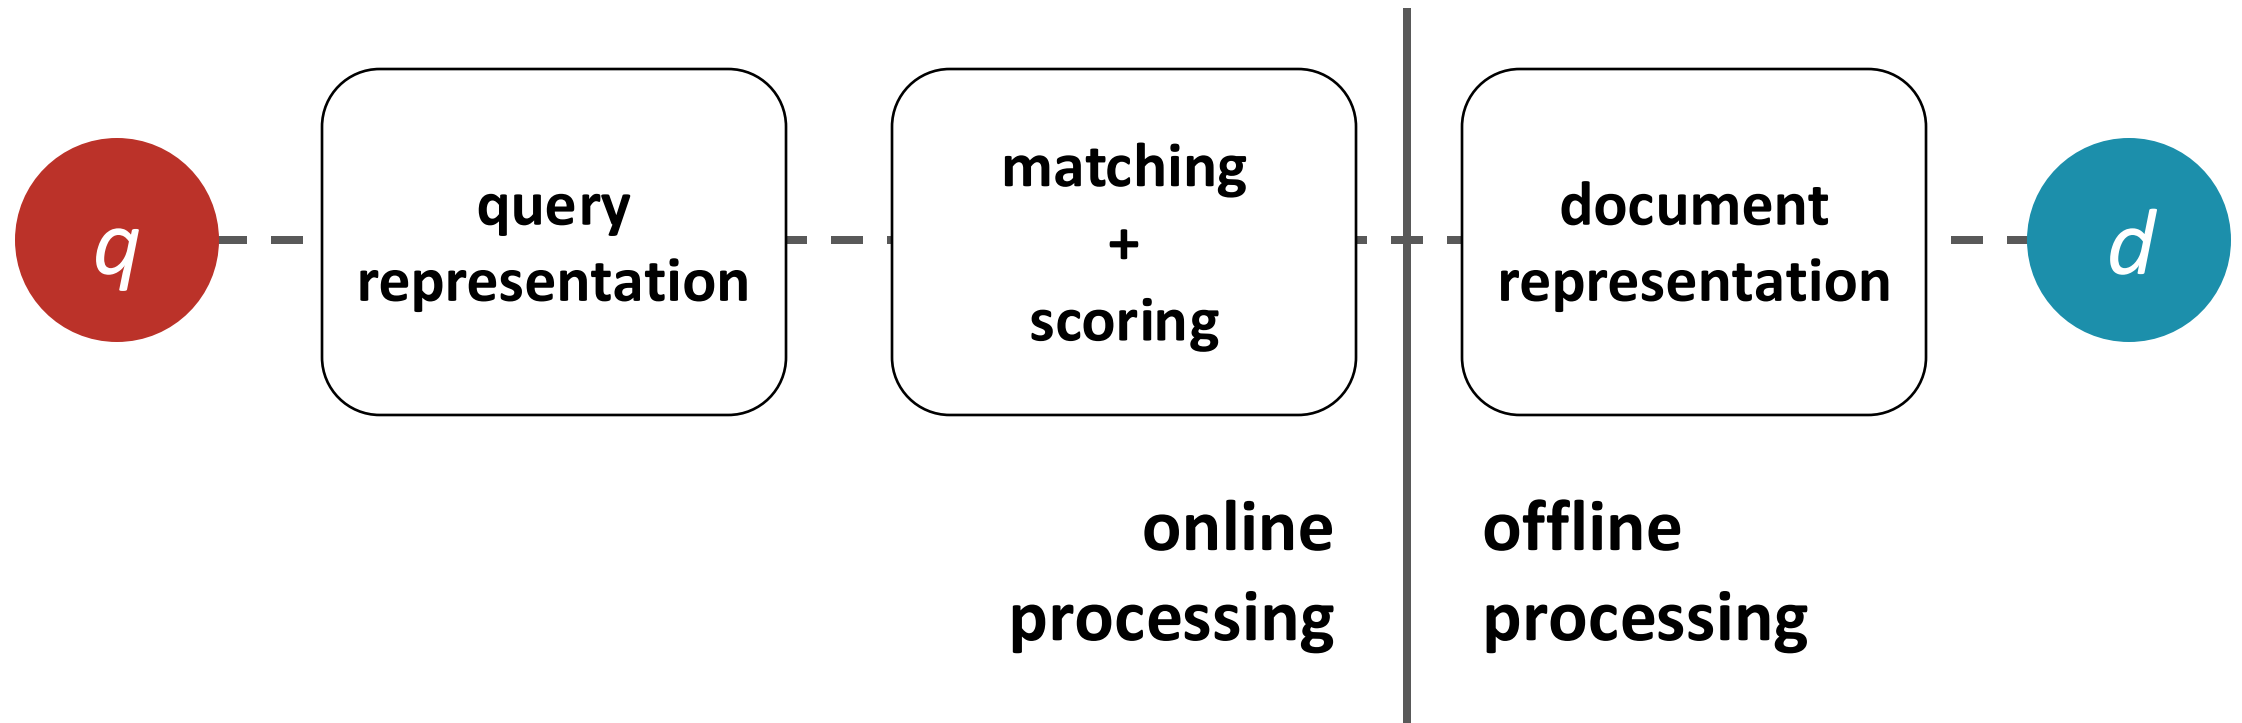
**What do
search
engineers
do?**

The search problem



$$f(q, d)$$

Search pipeline



(Continuous) offline processing

Document acquisition

Document understanding

Document indexing

Document acquisition

The Web is huge

- Trillions of known URLs, billions fetched

The Web is constantly evolving

- Updates, additions, deletions

Efficient crawling is key

- Must aim for coverage, but also freshness

Document understanding

Documents carry meaning

- Term-based matching as a first approximation
- Several techniques to leverage semantics

Documents vary in quality

- *Genuinely*: accessibility, readability, authority, depth
- *Maliciously*: content / link farms, misinformation

Document indexing

Efficient retrieval through indexes

- Like the index of a book
 - For each word, a list of documents it appears on
- Broken up into shards of millions of documents
 - 1000s of shards for the web index
- Plus per-document metadata
- Plus document embeddings

Online processing

Query understanding

Matching and scoring

Post-processing

Query understanding

Keywords are poor descriptions of the user's need

- Interaction and context also matter

Query understanding techniques can help

- Query segmentation, query scoping
- Query relaxation, query expansion
- Query embedding

Query understanding

Query scoping through semantic annotation

- [**san jose** convention center]
- [**matt cutts**]

Query expansion through acronym expansion

- [**gm** trucks] → [**general motors** trucks]
- [**gm** corn] → [**genetically modified** corn]

Matching and scoring

Send the query to all the shards

Each shard

- Finds matching documents
- Scores each query-document pair
- Sends back the top n documents

Combine all the top documents and sort by score

Ranking evaluation

Relevance is a user's prerogative

- We can observe changes in user behavior
- Or directly ask the user how we're doing

Evaluation is an empirical science

- It must be scientifically rigorous
- It must be economically viable

Course goals

Provide an introductory account of methods for building and evaluating search engines

Provide an exploration of recent advances and current research directions in the field

Course scope

System view

- Crawling, indexing, retrieval

Modeling view

- Ranking models

Behavioral view

- Ranking evaluation

Out-of-scope

We have dedicated courses for:

- Recommender systems
- Natural language processing
- Machine learning
- Data mining

Course grading (tentative)

Exams: 50%

Assignments: 40%

Seminars: 10%

Course attendance

“

O que é necessário para ser aprovado em uma dada atividade acadêmica curricular?

É necessário obter nota final igual ou superior a 60, em uma escala de 0 a 100, bem como a indicação de assiduidade, a qual deve ser igual ou superior a 75% (art. 12 das NGG).

Course materials: textbooks

[Search Engines: Information Retrieval in Practice](#)

by B. Croft, D. Metzler, and T. Strohman

[Introduction to Information Retrieval](#)

by C. Manning, P. Raghavan, and H. Schütze

[Modern Information Retrieval](#)

by R. Baeza-Yates and B. Ribeiro-Neto

Course materials: textbooks

[Information Retrieval: Implementing and Evaluating Search Engines](#)

by S. Büttcher, C. Clarke, and G. Cormack

[Text Data Management: A Practical Introduction to Information Retrieval and Text Mining](#)

by C. Zhai and S. Massung

Course materials: surveys

[Foundations and Trends in Information Retrieval](#)

by several authors

[Synthesis Lectures on Information Concepts, Retrieval, and Services](#)

by several authors

Other relevant material

General background

- Algorithms and data structures
- Basic statistics
- Basic linear algebra

Advanced readings

- [Google Scholar](#) is your friend

References

[Search Engines: Information Retrieval in Practice](#), Ch. 1

Croft et al., 2009

[How Google Works: A Google Ranking Engineer's Story](#)

Haahr, SMX West 2016

[Ten blue links on Mars](#)

Clarke et al., WWW 2017

Pre-course survey

Fill in a short survey describing your past experience and expectations related to the course

- <https://forms.gle/7mcatGc5LtAFM2ta7>





UNIVERSIDADE*FEDERAL
DE*MINAS*GERAIS

Coming next...

Search Architecture

Rodrygo L. T. Santos
rodrygo@dcc.ufmg.br