




UNIVERSIDADE\*FEDERAL  
DE\*MINAS\*GERAIS

Information Retrieval

# Experimental Methods

Rodrygo L. T. Santos  
[rodrygo@dcc.ufmg.br](mailto:rodrygo@dcc.ufmg.br)

# One problem



$$f(q, d)$$

# Many solutions

Similarity-based models

Probabilistic models

Extended models

Machine-learned models

# Why evaluate?

Lots of alternative solutions

- Which one to choose?
- How to improve upon them?

Evaluation enables an informed choice

- Rigor of science
- Efficiency of practice

# What to evaluate?

Three fundamental types of IR research

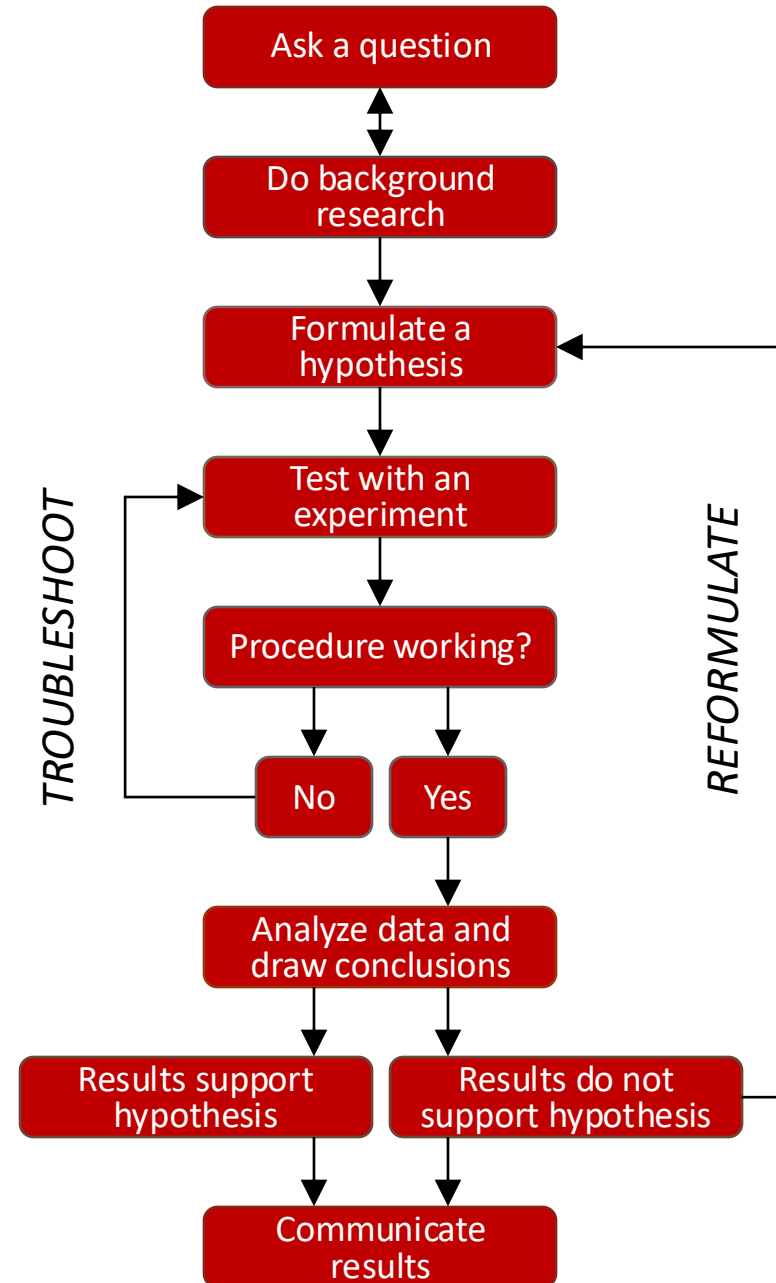
- Systems (efficiency)
- Methods (effectiveness)
- Applications (user utility)

Evaluation plays a critical role for all three

- Our primary focus is on “methods” research

# How to evaluate?

Scientifically, of course!



# Asking questions

What problem are you trying to solve?

- Or in IR parlance, what task?

Hard to solve an ill-defined task!

- Is it a well-known task? Review the literature!
- Is it unlike anything done before?

# Asking (new) questions

Characterize the task

- How is the system used?
- What are the inputs? Outputs?
- How do you define success?



# Formulating hypotheses

A hypothesis must be falsifiable

- Ideally concerning an isolated component  
*e.g., length normalization improves ranking*

It either holds or does not...

- ... with respect to the considered data (scope)
- ... perhaps under certain conditions (extent)

# Performing experiments

## Key components

- Experimental setup
- Analysis of results

## Key concern: ***reproducibility***

- Must specify each and every detail needed for reproducing our method and the experiment

# Experimental setup

Research questions

Evaluation methodology

Evaluation benchmarks

Reference comparisons

Parameter tuning

# Research questions

Methods are not devised arbitrarily

- We always have a hypothesis (whether implicit or explicit) for why our work should improve
- Even the best results are useless if nobody understands what you are trying to solve

So, spell out your research questions!

# Evaluation methodology

We want to know

- What users consider relevant

We can observe

- What users tell us (explicit feedback)
- What users do (implicit feedback)

These are *noisy* measurements

# Evaluation methodology

Prospective experiments

- How well can we predict future preferences?

Benchmarked using live user interactions

- Poorly reproducible
- Highly realistic

# Evaluation methodology

Retrospective experiments

- How well can we predict (hidden) past preferences?

Benchmarked using static test collections

- Highly reproducible
- Poorly realistic

# Evaluation methodology

## Feedback

- Implicit
- Explicit

## Mode

- Retrospective
- Prospective

	<i>retrospective</i>	<i>prospective</i>
<i>implicit</i>	counterfactual evaluation	online evaluation
<i>explicit</i>	offline evaluation	



# Public test collections

## Text REtrieval Conference

- TREC has collections on Web, blog, tweet, video, question-answering, legal documents, medical records, chemicals, genomics, ... search

<http://trec.nist.gov/tracks.html>

<http://trec.nist.gov/data.html>

# You can build your own

Three core components

- A corpus of documents
- A set of users' queries
- A map of users' relevance assessments

# You can build your own

Document corpus

- Go crawl it!

Queries

- The more the better (e.g., at least 50)
- Representative of the population (e.g., from a log)

Relevance judgments

# How to judge relevance?

Who does it?

- Hired judges? Volunteers? Experts? Live users?

What are the instructions?

- Short queries? Long narratives?

What is the level of agreement?

- Redundancy to counter subjectivity

# What to judge for relevance?

Exhaustive assessment is not practical

- Alternative: document sampling

Stratified sampling via pooling

- Top  $k$  results from  $m$  rankers merged
- Unique (up to  $km$ ) results submitted for judgment

Generally robust for evaluating new rankers

# Reference comparisons (aka baselines)

*My method achieves 0.9 precision*

- Meaningless without a reference comparison
- Rephrasing: is it better or worse?

Choice of baseline depends on the hypothesis

- Key question: what are you trying to show?

# Choosing baselines

## Vanilla baselines

- Have the proposed effect turned off  
*e.g., ranking without length normalization*

## Competing baselines

- Exploit the proposed effect in a different manner  
*e.g., alternative length normalization*

# Parameter tuning

Your method may have parameters

- Your baselines may also have parameters  
e.g.,  $b$  for pivoted length normalization

Which parameters need tuning?

- Which can stay fixed?
- How to tune?



# Analysis of results

Measure, compare, slice and dice results

- Helps prove (or disprove) your hypotheses
- Demonstrates how your methods or systems compare against the existing state-of-the-art
- Provides fundamental insights into the underlying research problems being addressed

# Evaluation metrics

General form:  $\Delta(R, G)$

- $R$ : ranking produced by model  $f$  for query  $q$
- $G$ : ground-truth produced for query  $q$

Metrics should be chosen according to the task

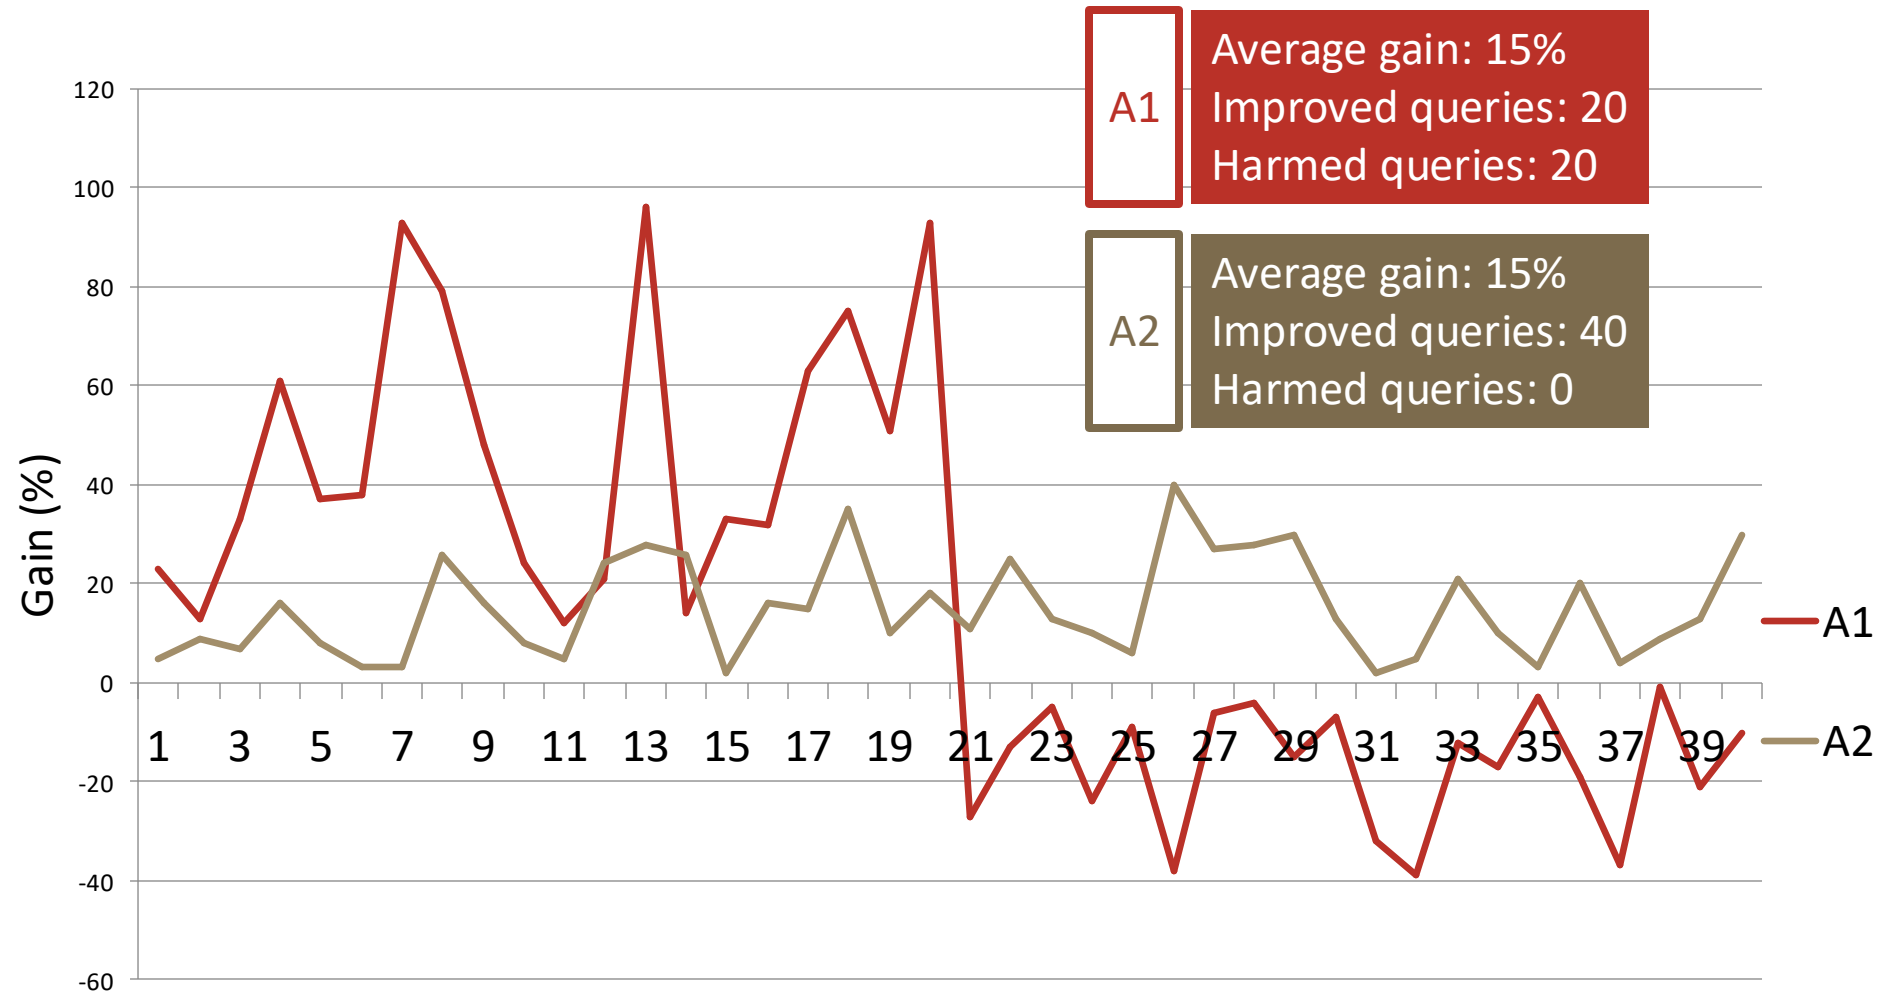
- Web search (precision) vs. legal search (recall)  
(more on next class)

# Results significance

Effectiveness varies across queries

- Large average improvement may not be consistent
- Might improve a lot on some queries, hurt on many

# Variable effectiveness



# Results significance

Effectiveness varies across queries

- Large average improvement may not be consistent
- Might improve a lot on some queries, hurt on many

Improvements should be tested for significance

- Statistical significance (see next class)
- Practical significance

# Deeper analyses

*My method beats the baseline...*

- *... phew, let's call it a victory and go home! #NOT*

Deeper analyses may provide further insights

- Why the method works
- When the method works
- And when it doesn't!

# Deeper analyses

Parameter sensitivity analysis

- How sensitive is the method to its parameters?

Breakdown analysis

- How does it perform for different queries?

Failure analysis

- What are the main reasons for failure?

# Summary

Experimentation drives search innovation

- Experiments should be economically practical
- Experiments should be scientifically rigorous
- Experiments should be reproducible
- Experiments should provide insights



# References

[Experimental methods for information retrieval](#)

Metzler and Kurland, SIGIR 2012

[Introduction to Information Retrieval](#), Ch. 8

Manning et al., 2008

[Search Engines: Information Retrieval in Practice](#), Ch. 8

Croft et al., 2009

Coming next...

# Offline Evaluation

Rodrygo L. T. Santos  
rodrygo@dcc.ufmg.br