


Information Retrieval

Quality Models

Rodrygo L. T. Santos
rodrygo@dcc.ufmg.br

The ranking problem



$$f(q, d)$$


Query likelihood model

$$\begin{aligned} f(q, d) &\approx P(q, d) \\ &= P(q|d)P(d) \quad \text{Bayes' rule} \end{aligned}$$

Two core components


- $P(q|d)$: query likelihood
- $P(d)$: document prior

The ranking problem



$$f(q, d)$$

The ranking problem



$f(d)$

Quality as prior relevance

High quality of the web document content increases the a priori probability of the document being relevant

- a.k.a. document prior

Quality factors should be combined in a way that directly improves the retrieval effectiveness

- e.g., nDCG or MAP

Document quality in web search

Web is decentralized and heterogeneous

- Different authority
- Different goals
- Different credibility
- Different publishing standards

Document quality in web search

Quality of a web page is determined by many factors

- Original, up-to-date content of genuine value
- Links to related resources
- Layout for easy reading and navigation

Continuous spectrum from high-quality pages to spam

- Most web documents are somewhere in between

Document quality in web search

“

*As pure web spam has decreased over time,
attention has shifted instead to sites with
shallow or low-quality content.*

◦ Matt Cutts, 2011

Document quality in web search

Document quality in search engines

- Promote high-quality content
- Demote low-quality content

Two broad sources of quality evidence

- On-document evidence
- Off-document evidence

On-document evidence

Verbosity

Readability

Cohesiveness

Navigability

Support

Verbosity

Document nominal length (Singhal et al., 1996)

- Full length in tokens

Document visible length (Zhu & Gauch, 2000)

- Content actually rendered

Document title length (Bendersky et al., 2011)

- Measures descriptiveness of page metadata

Readability

Average term length (Kanungo & Orr, 2009)

- Longer terms denote thoughtful selection

Stopwords ratio/coverage (Ntoulas et al., 2006)

- Correlated with informativeness

Fraction of table text (Bendersky et al., 2011)

- High fraction denotes poor readability

Cohesiveness

Entropy of the page content (Bendersky et al., 2011)

- Lower entropy denotes better focus

$$H(\theta_d) = - \sum_{t \in d} P(t|d) \log P(t|d)$$

Navigability

URL length and depth (Kraaij et al., 2002)

- Lower length and depth denote easier navigation

URL type (domain, subdomain, path, file)

- Homepages tend to be of type domain

Support

Fraction of anchor text (Ntoulas et al., 2006)

- Reasonable amount conveys factuality
- Excess denotes shallowness

Off-document evidence

On-document evidence provides valuable evidence about the quality of the document

- Also prone to manipulation by the document author

Off-document evidence isn't immune, but is less biased

- Hyperlink structure
- Click-through data

Link analysis

Links are a key component of the Web

- Important for navigation, but also for search

Two complementary sources of information

`Example website`

- “Example website”: anchor text
- “http://example.com”: destination link

Authority

Billions of web pages, more or less informative

- Links can be viewed as information about the popularity (authority?) of a web page

Inlink count could be used as a simple measure

- Link analysis algorithms like PageRank provide more reliable ratings (less susceptible to link spam)

Random surfer model

Choose a random number r between 0 and 1

- If $r < \lambda$: go to a random page (avoid getting stuck)
- If $r \geq \lambda$: follow a random link from the current page

Repeat


PageRank

PageRank of page u is the probability that the “random surfer” will be looking at u as $t \rightarrow \infty$

$$PR^{(t+1)}(u) = \frac{\lambda}{n} + (1 - \lambda) \sum_{v \in I_u} \frac{PR^{(t)}(v)}{|O_v|}$$

- I_u : inlinks of page u
- O_v : outlinks of page v

PageRank example ($\lambda = 0$)



$$PR^{(t+1)}(A) = \frac{PR^{(t)}(C)}{1}$$

$$PR^{(t+1)}(B) = \frac{PR^{(t)}(A)}{2}$$

$$PR^{(t+1)}(C) = \frac{PR^{(t)}(A)}{2} + \frac{PR^{(t)}(B)}{1}$$

	$PR^{(t)}(A)$	$PR^{(t)}(B)$	$PR^{(t)}(C)$
$t = 0$	0.33	0.33	0.33
$t = 1$	0.33	0.17	0.50
$t = 2$	0.50	0.17	0.33
$t = 3$	0.33	0.25	0.42
...			
$t = \infty$	0.40	0.20	0.40

Convergence check

Typical stopping criteria

- $\|\overrightarrow{PR}^{(t)} - \overrightarrow{PR}^{(t-1)}\| < \epsilon$

Setting ϵ

- Small ϵ : slow convergence, accurate PR
- Large ϵ : fast convergence, inaccurate PR

Click-through rate

Global click likelihood

$$CTR(d) = \frac{\sum_{q \in L} 1_{K_q}(d)}{\sum_{q \in L} 1_{R_q}(d)}$$

- $q \in L$: a query in the log
- K_q : documents clicked for q
- R_q : documents displayed for q

What is web spam?

Spamming = any deliberate action solely in order to boost a web page's position in search engine results, incommensurate with that page's real value

- Spam = web pages that are the result of spamming

Approximately 10-15% of web pages are spam

- High premium to appear on the first page of results

Web spam taxonomy

Boosting techniques

- Aim at scoring high for topicality / authoritativeness

Hiding techniques

- Aim at hiding the use of boosting

Boosting techniques

Term spamming

- Inflate content to appear relevant to many queries

Link spamming

- Creating link structures that boost authority

Term spamming

Repetition

- Repeat one or a few specific terms (e.g., free, cheap)
- Goal is to subvert tf-idf ranking schemes

Dumping

- Add a large number of unrelated terms
- Goal is to match a variety of queries

Term spamming

Weaving

- Randomly stick spam terms along legitimate content

Phrase stitching

- Glue together sentences from different sources

Link spamming

Spammer's inaccessible pages


- Most of the Web

Spammer's accessible pages

- e.g., blog comments pages

Spammer's own pages

- Completely controlled by spammer



Link farms

Spammer's goal

- Maximize the authority of target page t

Spammer's approach

- Get many links from accessible pages to page t
- Construct “link farm” to get multiplier effect

Hiding techniques

Content hiding

- Use same color for text and page background

Cloaking

- Return different page to crawlers and browsers

Detecting spam

Term spamming

- Analyze text using statistical classifiers
- Also useful: near duplicate detection

Link spamming

- Trust propagation (ham pages link to ham pages)
- Open research area

**How to
combine
relevance
and
quality?**

Quality as a static score

Query-independent scoring


- Typically computed offline

Precomputed scores leveraged in multiple ways

- As a multiplier in vector space models
- As a prior in probabilistic models
- As a feature in feature-based models

How effective is it?

(Bendersky et al., 2011)



Summary

Document quality can be heterogeneous

- Quality models can help distinguish between documents with similar relevance scores
- Also useful for queries that explicitly seek for high-quality content (authority, readability)
- And to combat adversarial behavior (spam, fakes)

References

[Quality-biased ranking of web documents](#)

Bendersky et al., WSDM 2011

[The importance of prior probabilities for entry page search](#)

Kraaij et al., SIGIR 2002

[The PageRank citation ranking: bringing order to the Web](#)

Page and Brin, Tech report 1999

References

[Relevance weighting for query independent evidence](#)

Craswell et al., SIGIR 2005

[Adversarial web search](#)

Castillo and Davison, FnTIR 2011

[Fake news detection on social media](#)

Shu et al., SIGKDD Expl. 2017

Coming next...

Feedback Models

Rodrygo L. T. Santos
rodrygo@dcc.ufmg.br