


Information Retrieval

Learning to Rank: Pointwise

Rodrygo L. T. Santos


rodrygo@dcc.ufmg.br

The ranking problem



$$f(q, d)$$

Learning to rank



$f(\mathbf{x})$

Learning to rank

Feature-based representation

- Individual models as ranking “features”

Discriminative learning

- Effective models learned from data
- Aka machine-learned ranking

Building blocks

Goal is to learn a ranking model

- $f : \mathcal{X} \rightarrow \mathcal{Y}$

That minimizes some loss function

- $\mathcal{L} : f(\mathcal{X}) \times \mathcal{Y} \rightarrow \mathcal{R}$

Ideally, we would like a low test error

- We settle for a good training error / capacity trade-off

What is a learning algorithm?

Given

- A family of functions \mathcal{F} (e.g., linear, trees, neural nets)
- A measure of loss \mathcal{L} (error + capacity)

Learning can be cast as an optimization problem

- $f^* = \operatorname{argmin}_{f \in \mathcal{F}} \mathcal{L}(f, \text{train})$
 $= \operatorname{argmin}_{f \in \mathcal{F}} \operatorname{Err}(f, \text{train}) + \lambda \operatorname{Reg}(f)$

Classical algorithms

Linear learning algorithms

- Regression (LASSO, Ridge reg.)
- Classification (logistic reg., linear SVM, AdaBoost)

Non-linear learning algorithms

- Neural networks
- (Boosted) regression trees

Linear learning algorithms

\mathcal{F} is the set of linear functions

- $f_w(x) = w^T x$, where $x \in \mathcal{R}^d$, $w \in \mathcal{R}^d$
 - x : d -dimens. feature vector describing an example
 - w : weight vector defining the function f_w
 - $f_w(x)$ is the prediction of f_w for example x

Linear learning algorithms

Most typical regularizers

- L2 regularization, $\|w\|_2^2 = \sum_j w_j^2$
(prefer "flat" weights)
- L1 regularization, $\|w\|_1 = \sum_j |w_j|$
(prefer "sparse" weights)

Classical algorithms

Linear learning algorithms

- **Regression (LASSO, Ridge reg.)**
- Classification (logistic reg., linear SVM, AdaBoost)

Non-linear learning algorithms

- Neural networks
- (Boosted) regression trees

Linear learning for regression

Ordinary least squares

- $\mathcal{L}(w) = \frac{1}{2} \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)})^2$

Ridge regression

- Add L2 regularization ($+\lambda \|w\|_2^2$)

LASSO

- Add L1 regularization ($+\lambda \|w\|_1$)

Linear learning for regression

Ordinary least squares

- $\mathcal{L}(w) = \frac{1}{2} \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)})^2$

Typical optimization strategy

- Compute the gradient of $\mathcal{L}(\cdot)$ as $\nabla \mathcal{L}(\cdot)$
- Take a step (iterate) in the opposite direction
 $w = w - \alpha \nabla \mathcal{L}(w)$, where α is the learning rate

Computing gradients

Gradient as a d -dimens. vector of partial derivatives

- $\nabla \mathcal{L}(w) = \left(\frac{\partial \mathcal{L}(w)}{\partial w_1}, \frac{\partial \mathcal{L}(w)}{\partial w_2}, \dots, \frac{\partial \mathcal{L}(w)}{\partial w_d} \right)$

How to compute each partial derivative?

- Manual derivation... d may be very large!
- Automatically

Computing gradients

Ordinary least squares

- $\mathcal{L}(w) = \frac{1}{2} \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)})^2$

For a given dimension k

- $\frac{\partial}{\partial w_k} \mathcal{L}(w) = \frac{\partial}{\partial w_k} \frac{1}{2} \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)})^2$

Computing gradients

$$\begin{aligned}\frac{\partial}{\partial w_k} \mathcal{L}(w) &= \frac{\partial}{\partial w_k} \frac{1}{2} \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)})^2 \\ &= \sum_{i=1}^m \frac{\partial}{\partial w_k} \frac{1}{2} (f_w(x^{(i)}) - y^{(i)})^2 \\ &= \sum_{i=1}^m 2 \frac{1}{2} (f_w(x^{(i)}) - y^{(i)}) \frac{\partial}{\partial w_k} (f_w(x^{(i)}) - y^{(i)}) \\ &= \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)}) \frac{\partial}{\partial w_k} f_w(x^{(i)})\end{aligned}$$

Computing gradients

$$\begin{aligned}\frac{\partial}{\partial w_k} f_w(x^{(i)}) &= \frac{\partial}{\partial w_k} w^T x^{(i)} \\ &= \frac{\partial}{\partial w_k} \sum_{j=1}^d w_j x_j^{(i)} \\ &= \frac{\partial}{\partial w_k} \left(w_k x_k^{(i)} + \sum_{j \neq k} w_j x_j^{(i)} \right) \\ &= x_k^{(i)}\end{aligned}$$

Computing gradients

$$\begin{aligned}\frac{\partial}{\partial w_k} \mathcal{L}(w) &= \sum_{i=1}^m \left(f_w(x^{(i)}) - y^{(i)} \right) \frac{\partial}{\partial w_k} f_w(x^{(i)}) \\ &= \sum_{i=1}^m \underbrace{\left(f_w(x^{(i)}) - y^{(i)} \right)}_{\text{prediction error}} \underbrace{x_k^{(i)}}_{k^{\text{th}} \text{ feature score}}\end{aligned}$$

Computing gradients

$$\begin{aligned}\frac{\partial}{\partial w_k} \mathcal{L}(w) &= \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)}) \frac{\partial}{\partial w_k} f_w(x^{(i)}) \\ &= \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)}) x_k^{(i)}\end{aligned}$$

- $m = |D|$: batch gradient descent
- $m = c \ll |D|$: mini-batch gradient descent
- $m = 1$: stochastic gradient descent

Classical algorithms

Linear learning algorithms

- Regression (LASSO, Ridge reg.)
- **Classification (logistic reg., linear SVM, AdaBoost)**

Non-linear learning algorithms

- Neural networks
- (Boosted) regression trees

Linear learning for classification

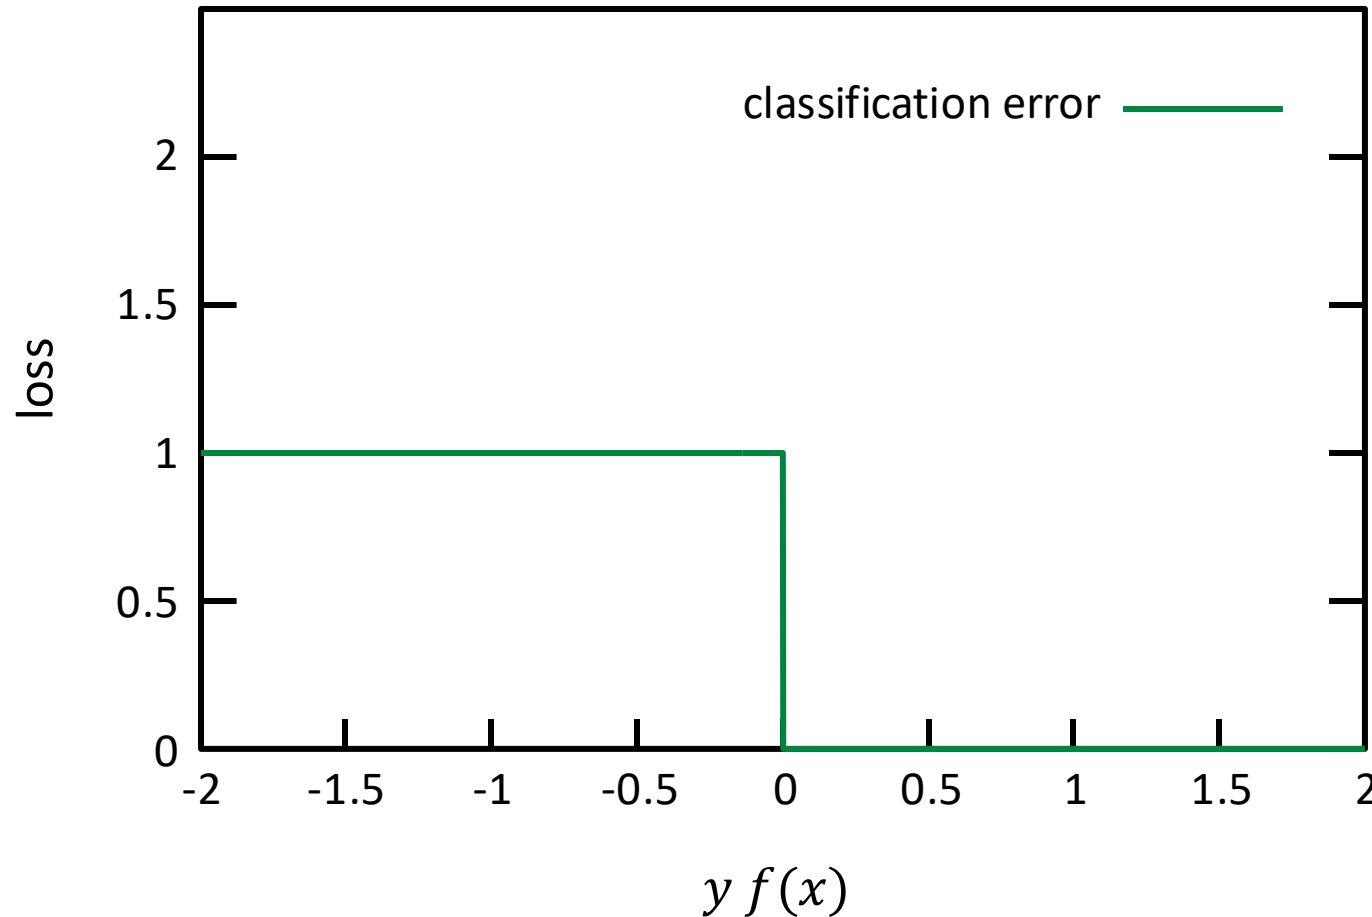
Prediction

- $f_w(x) = w^T x$
 - > 0 predicts $+1$
 - ≤ 0 predicts -1

Classification error


- $\mathcal{L}(w) = \frac{1}{m} \sum_{i=1}^m \mathbf{1}(y^{(i)} f_w(x^{(i)}) < 0)$
- **Problem:** non-differentiable!

Upper-bounding classification error



Classification error
 $\mathbf{1}(y f_w(x) < 0)$


Upper-bounding classification error



Logistic loss

$$\log(1 + \exp(-y f_w(x)))$$


Upper-bounding classification error



Hinge loss

$$\max(0, 1 - y f_w(x))$$

Upper-bounding classification error



Exponential loss
 $\exp(-y f_w(x))$

Linear learning for classification

Logistic regression

- Logistic loss + L1 or L2 regularizer

Support Vector Machines (SVM)

- Hinge loss + L2 regularizer

AdaBoost

- Exponential loss + L1 regularizer

Classical algorithms

Linear learning algorithms

- Regression (LASSO, Ridge reg.)
- Classification (logistic reg., linear SVM, AdaBoost)

Non-linear learning algorithms

- **Neural networks**
- (Boosted) regression trees


Neural networks

Composable functions through several layers

- Inspired by biological neural networks (your brain)


Each layer performs a linear operation potentially followed by a non-linear activation

- $\ell_0 = x$ and $\ell_i = \underbrace{\sigma_i}_{\text{activation}}(\underbrace{W_i \ell_{i-1}}_{\text{weights}} + \underbrace{b_i}_{\text{biases}})$



sigmoid function

Example: shallow (2-layer) network



$$\ell_0 = x$$

$[d \times 1]$


$$\ell_1 = \sigma_1(W_1 \ell_0 + b_1)$$

$[h_1 \times 1]$ $[h_1 \times d]$ $[d \times 1]$ $[h_1 \times 1]$

$$\ell_2 = \sigma_2(W_2 \ell_1 + b_2)$$

$[1 \times 1]$ $[1 \times h_1]$ $[h_1 \times 1]$ $[1 \times 1]$

Example: deep (5-layer) network



$$\ell_0 = x$$

$[d \times 1]$

$$\ell_1 = \sigma_1(W_1 \ell_0 + b_1)$$

$[h_1 \times 1]$ $[h_1 \times d]$ $[d \times 1]$ $[h_1 \times 1]$

$$\ell_2 = \sigma_2(W_2 \ell_1 + b_2)$$

$[h_2 \times 1]$ $[h_2 \times h_1]$ $[h_1 \times 1]$ $[h_2 \times 1]$

$$\ell_3 = \sigma_3(W_3 \ell_2 + b_3)$$

$[h_3 \times 1]$ $[h_3 \times h_2]$ $[h_2 \times 1]$ $[h_3 \times 1]$

$$\ell_4 = \sigma_4(W_4 \ell_3 + b_4)$$

$[h_4 \times 1]$ $[h_4 \times h_3]$ $[h_3 \times 1]$ $[h_4 \times 1]$

$$\ell_5 = \sigma_5(W_5 \ell_4 + b_5)$$

$[1 \times 1]$ $[1 \times h_4]$ $[h_4 \times 1]$ $[1 \times 1]$

Neural network learning

Parameters (W_i 's and b_i 's) learned via gradient descent

- Find $\{W_i, b_i\}$ which minimize loss

Works with any differentiable loss

- Cross-entropy commonly used
- L2 regularization on $\{W_i, b_i\}$

Architecture selected using validation data

Classical algorithms


Linear learning algorithms

- Regression (LASSO, Ridge reg.)
- Classification (logistic reg., linear SVM, AdaBoost)

Non-linear learning algorithms

- Neural networks
- **(Boosted) regression trees**

Regression trees



$x = (0.3, 0.6, 0.2, 0.5)$

$x_2 > 0.5$? **T**

$x_4 > 0.8$? **F**

$\therefore f(x) = 0.8$

Examples travel the tree from the root to one leaf

Each node performs a test on the input x

Each leaf corresponds to a prediction $f(x)$

Greedy learning

1. Start with a tree containing only the root
2. Splitting: for each node with depth $<$ max depth
 - Find best test and create two leaves from the node
 - Find best prediction for these two leaves
3. Rep. #2 until no more nodes with depth $<$ max depth


Greedy learning

$$f(x) = 0.5$$

Best split: $x_2 > 0.5$

Best predictions: $f(x) = 0.7$ if $x_2 > 0.5$
 $f(x) = 0.2$ otherwise


Greedy learning



Best split: $x_2 > 0.5$

Best predictions: $f(x) = 0.7$ if $x_2 > 0.5$
 $f(x) = 0.2$ otherwise


Greedy learning



Best split: $x_4 > 0.8$

Best predictions: $f(x) = 0.5$ if $x_4 > 0.8$
 $f(x) = 0.8$ otherwise


Greedy learning



Best split: $x_4 > 0.8$

Best predictions: $f(x) = 0.5$ if $x_4 > 0.8$
 $f(x) = 0.8$ otherwise

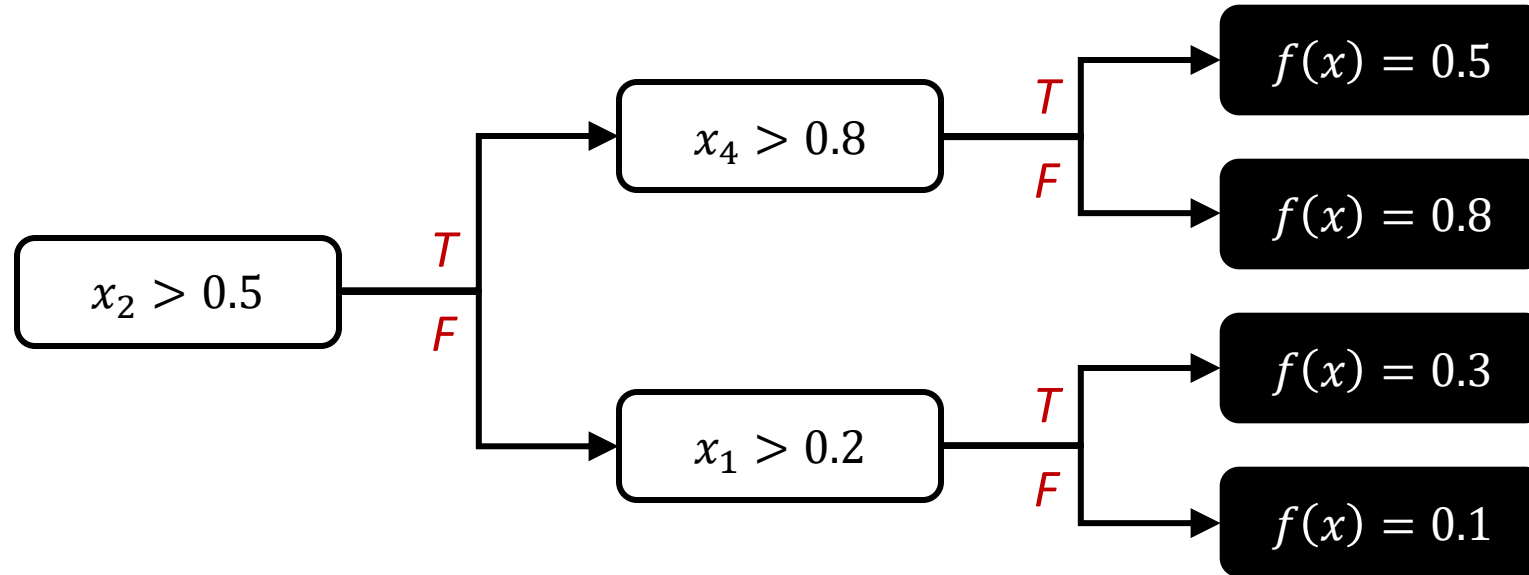
Greedy learning



Best split: $x_1 > 0.2$

Best predictions: $f(x) = 0.3$ if $x_1 > 0.2$
 $f(x) = 0.1$ otherwise

Greedy learning



Best split: $x_1 > 0.2$

Best predictions: $f(x) = 0.3$ if $x_1 > 0.2$

$f(x) = 0.1$ otherwise

Greedy learning a regression tree

1. Start with a tree containing only the root
2. Splitting: for each node with depth $<$ max depth
 - Find best test and create two leaves from the node
 - Find best prediction for these two leaves
3. Rep. #2 until no more nodes with depth $<$ max depth

Splitting a node

Given all S_k examples reaching node k , find:

- i, τ : defining the test $x_i > \tau$

Let $S_k^l \equiv S_k: x_i \leq \tau$ and $S_k^r \equiv S_k: x_i > \tau$

Minimize: $\mathcal{L} = \sum_{(x,y) \in S_k^l} (\bar{y}^l - y)^2 + \sum_{(x,y) \in S_k^r} (\bar{y}^r - y)^2$

$$\bullet \bar{y}^l \equiv \frac{1}{|S_k^l|} \sum_{(x,y) \in S_k^l} y \quad \text{and} \quad \bar{y}^r \equiv \frac{1}{|S_k^r|} \sum_{(x,y) \in S_k^r} y$$

$f(x)$ in the left branch

$f(x)$ in the right branch

Regression trees

Advantages

- Low computation cost for prediction
- Easy to interpret

Disadvantages

- Only models piecewise constant functions
- Learning deep trees requires lots of examples

Gradient boosted regression trees (GBRTs)

A GBRT is an ensemble of regression trees

- $f(x) = \sum_{t=1}^T h_t(x)$, where h_t is a regression tree

Advantages

- Models more complex functions than a single tree
- Learning many shallow trees requires less training data than learning one deep tree

Gradient boosting

For $t = 1, \dots, T$

- Current model: $f_t = \sum_{i=1}^{t-1} h_i$
- Learn h_t^* to "correct" f_t
 - $f_{t+1} = f_t + h_t^* \approx y$
 $\therefore h_t^* \approx \underline{y - f_t} \quad \text{residual}$
 - $h_t^* = \operatorname{argmin}_{h_t} \sum_{(x,y)} (h_t(x) - (y - f_t(x)))^2$
 $= \operatorname{argmin}_{h_t} \sum_{(x,y)} (h_t(x) - (-\alpha \nabla \mathcal{L}(f_t)))^2$

Summary

Pointwise approaches borrowed from other tasks

- Regression, classification, ordinal classification

Straightforward, yet not quite suited for ranking

- Ranking requires getting relative scores right
- Higher positions matter more than lower positions
- All queries are equally important

References

[Learning to rank for information retrieval](#)

Liu, FnTIR 2009

[Learning to rank for information retrieval](#)

Liu, 2011

[Learning to rank for information retrieval and natural language processing](#)

Li, 2014

References

[Learning for Web rankings](#)

Grangier and Paiement, 2011

[Stanford's CS229 lecture notes](#)

Ng, 2017

[TensorFlow and deep learning without a PhD](#)

Gorner, 2017

References

[Mathematics for machine learning](#)

Deisenroth et al., 2018

Coming next...

Learning to Rank: Pairwise and Listwise

Rodrygo L. T. Santos
rodrygo@dcc.ufmg.br