


Information Retrieval

Learning to Rank: Pairwise and Listwise

Rodrygo L. T. Santos


rodrygo@dcc.ufmg.br

The ranking problem



$$f(q, d)$$

Learning to rank



$f(\mathbf{x})$

Learning to rank

Feature-based representation

- Individual models as ranking “features”

Discriminative learning

- Effective models learned from data
- Aka machine-learned ranking

Pointwise approach

Several approaches

- Regression-based
- Classification-based
- Ordinal regression-based



$$\{(x_1, y_1), (x_2, y_2), (x_3, y_3), \dots, (x_n, y_n)\}$$

Limitations of the pointwise approach

Ranking requires getting relative scores right

- Pointwise approaches learn absolute scores

Higher positions should matter more than lower ones

- Pointwise loss functions are agnostic to positions

Queries should be equally important

- Queries with many relevant documents dominate

Pairwise approach

Pairwise classification-based

- RankNet
- RankBoost
- Ranking SVM
- IR-SVM



$$\{(x_1, x_2, 1), (x_2, x_1, 0), \\ (x_1, x_3, 1), (x_3, x_1, 0), \\ \dots$$

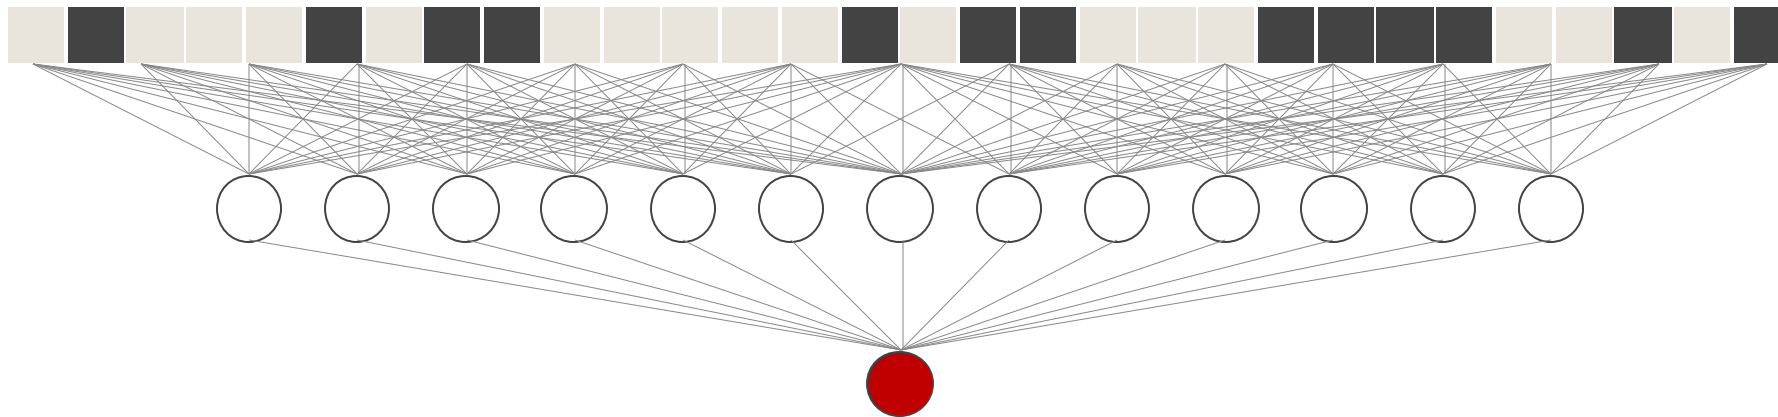
$$(x_{n-1}, x_n, 1), (x_n, x_{n-1}, 0)\}$$

RankNet

(Burges et al., ICML 2005)

Shallow (2-layer) neural network model

- Sigmoid activations



Gradient descent optimizer

RankNet

(Burges et al., ICML 2005)

Pairwise prediction converted to probability


$$\circ \hat{y}_{uv} = f(x_u, x_v) = \frac{\exp(f(x_u) - f(x_v))}{1 + \exp(f(x_u) - f(x_v))} \quad (\text{logistic function})$$

Cross entropy loss

$$\circ \mathcal{L}(f; x_u, x_v, y_{uv}) = -\overbrace{y_{uv}}^{\text{relevant}} \log \overbrace{\hat{y}_{uv}}^{\text{predicted relevant}} \\ - \underbrace{(1 - y_{uv})}_{\text{non-relevant}} \log \underbrace{(1 - \hat{y}_{uv})}_{\text{predicted non-relevant}}$$

Ranking SVM

(Herbrich et al., ALMC 2000; Joachims, KDD 2002)




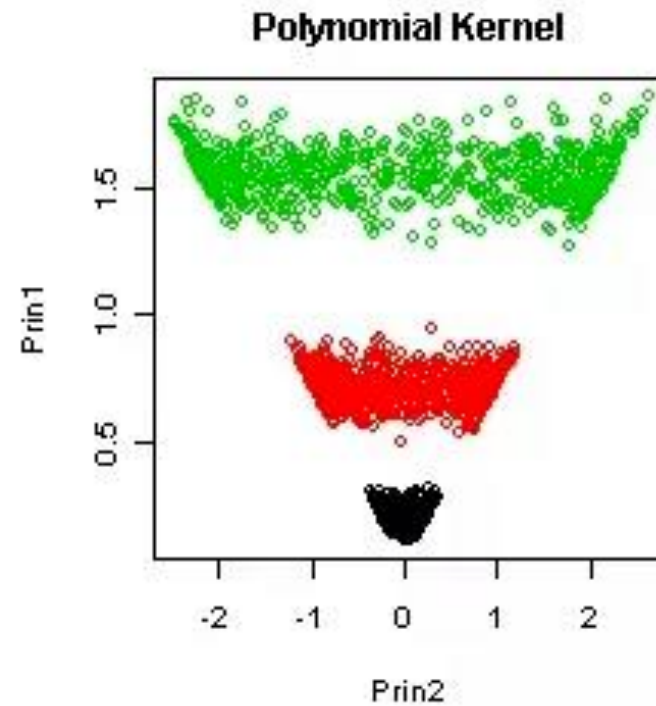

Hinge loss

- $\mathcal{L}(f; x_u, x_v, y_{uv}) = \max(0, 1 - y_{uv}(f(x_u) - f(x_v)))$

Nice properties inherited from standard SVM

- Good generalization via margin maximization
- Non-linear models via the kernel trick

Kernel trick



Limitations of the pairwise approach

Pairwise labels ignore graded relevance

- $(x_1, x_2, 1)$, regardless of the grades of x_1 and x_2

Query dominance exacerbated


- A query with more docs will have way more pairs

Ranking positions still not taken into account

- Swaps at the top more important than at the bottom

**Can we
optimize
ranking
metrics
directly?**

Ranking metrics generally non-differentiable



Piecewise constant functions

- **Flat:** zero derivative
- **Discontinuous:** undef derivative

LambdaRank

(Burges, NIPS 2006)


An extension of RankNet

- Ranking evaluation metrics (which are position-based) are directly used to define the gradient with respect to each document pair in the training process

Why is it feasible to directly define the gradient?

LambdaRank

(Burges, NIPS 2006)



$$\frac{\partial L}{\partial s_1} > \frac{\partial L}{\partial s_2}$$

LambdaRank

(Burges, NIPS 2006)

Gradient determines magnitude of updates

- $w = w - \alpha \nabla \mathcal{L}(w)$

Vector $\nabla \mathcal{L}(w) = \left(\frac{\partial \mathcal{L}(w)}{\partial w_1}, \frac{\partial \mathcal{L}(w)}{\partial w_2}, \dots, \frac{\partial \mathcal{L}(w)}{\partial w_d} \right)$

- $\frac{\partial}{\partial w_k} \mathcal{L}(w) = \sum_{\langle u,v \rangle} \lambda_{uv} x_k^{(i)}$

gradient magnitude *k^{th} feature score*
(prediction error)

LambdaRank

(Burges, NIPS 2006)

Lambda function

- An arbitrary surrogate for the gradient magnitude, assuming no particular loss function

$$\lambda_{uv} \equiv \frac{2^{y_u} - 2^{y_v}}{1 + \exp(f(x_u) - f(x_v))} |\Delta \text{nDCG}(x_u, x_v)|$$

LambdaMART

(Wu et al., Tech. Report 2008)

MART = Multiple Additive Regression Trees

- Commercial name for gradient boosted trees

Boosted tree version of LambdaRank

- Lambda functions guide the construction of weak learners (regression trees) via boosting
- $h_t^* = \operatorname{argmin}_{h_t} \sum_{(x,y)} (h_t(x) - (-\alpha \nabla \mathcal{L}(f_t)))^2$


Listwise approach

Metric-specific loss

- Optimize evaluation metrics

Non-metric-specific loss

- Optimize other listwise functions



Metric-specific listwise ranking

It is natural to directly optimize what is used to evaluate the ranking results, but not trivial

- Evaluation metrics such as nDCG and MAP are non-continuous and non-differentiable
- Most optimization techniques were developed to handle continuous and differentiable cases



SoftRank

(Taylor et al. LR4IR 2007/WSDM 2008)

Key idea: “soften” the evaluation metric

Score s_i as a random variable

- $P(s_i) = N(s_i | f(x_i), \sigma_s^2)$




SoftRank

(Taylor et al. LR4IR 2007/WSDM 2008)

We've constructed a score distribution per rank


- Non-deterministic scores make it possible to find a document in any ranking position

Map score distribution to rank distribution



SoftRank

(Taylor et al. LR4IR 2007/WSDM 2008)



Now each document has a non-deterministic position

- Each draw results in a permut.
- Each permut. has a given nDCG

Compute SoftNDCG as the expected nDCG over all possible permutations (smooth and differentiable)

Optimize via gradient descent

Metric-specific listwise ranking

Workarounds

- Soften or upper-bound non-smooth objective

Optimize the non-smooth objective directly

- Use genetic programming (RankGP)
- Use ranking evaluation metric to update the training data distribution via boosting (AdaRank)

AdaRank

(Xu and Li, SIGIR 2007)

Algorithm 2 Learning Algorithms for AdaRank

Input: document group for each query

Given: initial distribution \mathcal{D}_1 on input queries

For $t = 1, \dots, T$

Train weak ranker $f_t(\cdot)$ based on distribution \mathcal{D}_t .

Choose $\alpha_t = \frac{1}{2} \log \frac{\sum_{i=1}^n \mathcal{D}_t(i)(1+M(f_t, \mathbf{x}^{(i)}, \mathbf{y}^{(i)}))}{\sum_{i=1}^n \mathcal{D}_t(i)(1-M(f_t, \mathbf{x}^{(i)}, \mathbf{y}^{(i)}))}$

Update $\mathcal{D}_{t+1}(i) = \frac{\exp(-M(\sum_{s=1}^t \alpha_s f_s, \mathbf{x}^{(i)}, \mathbf{y}^{(i)}))}{\sum_{j=1}^n \exp(-M(\sum_{s=1}^t \alpha_s f_s, \mathbf{x}^{(j)}, \mathbf{y}^{(j)}))}$,

Output: $\sum_t \alpha_t f_t(\cdot)$.

| training set biased by
query difficulty

| single-feature weak ranker

| weighted accuracy of f_t

| update difficulty estimates

Non-metric-specific listwise ranking

Defining listwise loss functions based on the understanding of unique properties of ranking

Representative algorithms

- ListNet
- ListMLE
- BoltzRank

Ranking loss is non-trivial!

An example

- Function f : $f(A) = 4, f(B) = 2, f(C) = 3$ (ACB)
- Function h : $h(A) = 4, h(B) = 6, h(C) = 3$ (BAC)
- Ground-truth y : $y(A) = 6, y(B) = 4, y(C) = 3$ (ABC)

Which function (f or h) is closer to the ground-truth?

- According to nDCG, f should be closer to y !

ListNet

(Cao et al., ICML 2007)

Given f : $f(A) = 4, f(B) = 2, f(C) = 3$

- Deterministically: $P(ACB) = 1$

What if other permutations have a non-zero prob?

- For a k -permutation π^k under function f

$$P_f(\pi^k) = \frac{f(\pi_1)}{\sum_{i=1}^k f(\pi_i)} \times \frac{f(\pi_2)}{\sum_{i=2}^k f(\pi_i)} \times \dots \times \frac{f(\pi_{k-1})}{\sum_{i=k-1}^k f(\pi_i)}$$

ListNet

(Cao et al., ICML 2007)

Given f : $f(A) = 4, f(B) = 2, f(C) = 3$

- Deterministically: $P(ACB) = 1$

What if other permutations have a non-zero prob?


- Probabilistically: $P(ACB) = \frac{f(A)}{f(A)+f(C)+f(B)} \times$

$$\frac{f(C)}{f(C)+f(B)}$$

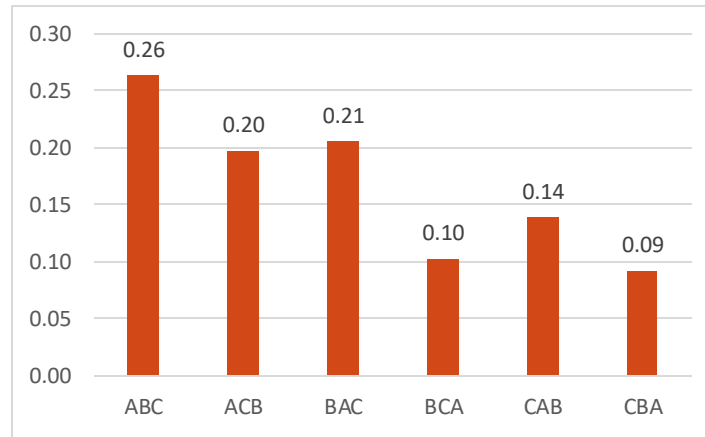
$$= \frac{4}{4+2+3} \times \frac{2}{2+3} = 0.27$$

Distance between ranked lists


$f: f(A) = 4, f(B) = 2, f(C) = 3$



$y: y(A) = 6, y(B) = 4, y(C) = 3$



$h: h(A) = 4, h(B) = 6, h(C) = 3$



$$D(P_y || P_f) = 0.10$$

$$D(P_y || P_h) = 0.14$$

f is closer!

KL divergence loss

Loss function = KL-divergence between two permutation probability distributions ($\varphi = \text{exp}$)

$$\circ \mathcal{L}(f; \mathbf{x}, y) = D(\underbrace{P_{\varphi(y)}}_{\text{ground-truth probability distro}} || \underbrace{P_{\varphi(f(\mathbf{x}))}}_{\text{predicted probability distro}})$$

Neural net model, gradient descent optimizer

Summary

Listwise approaches more closely model ranking

- Take all the documents associated with the same query as one learning instance
- Ranking position is visible to the loss function

State-of-the-art on standard benchmarks (LETOR)

- Particularly effective for top-heavy evaluation

RankLib tutorial

RankLib v2.18

- 0: MART (gradient boosted regression tree)
- 1: RankNet
- 2: RankBoost
- 3: AdaRank
- 4: Coordinate Ascent
- 6: LambdaMART
- 7: ListNet
- 8: Random Forests

RankLib tutorial

<https://sourceforge.net/p/lemur/wiki/RankLib%20How%20to%20use/>

```
> java -jar bin/RankLib-2.18.jar \  
    -train MQ2008/Fold1/train.txt \  
    -validate MQ2008/Fold1/vali.txt \  
    -test MQ2008/Fold1/test.txt \  
    -ranker 6 -metric2t NDCG@10 -metric2T ERR@10 \  
    -save mymodel.txt
```

References

[Learning to rank for information retrieval](#)

Liu, 2011

[Learning to rank for information retrieval and natural language processing](#)

Li, 2014

[Learning for Web rankings](#)

Grangier and Paiement, 2011

Coming next...

Neural Models

Rodrygo L. T. Santos
rodrygo@dcc.ufmg.br