

Pretrained Transformers for Text Ranking: BERT and Beyond

Andrew Yates, Rodrigo Nogueira, and Jimmy Lin

[@andrewyates](https://twitter.com/andrewyates)

[@rodrignogueira](https://twitter.com/rodrignogueira)

[@lintool](https://twitter.com/lintool)



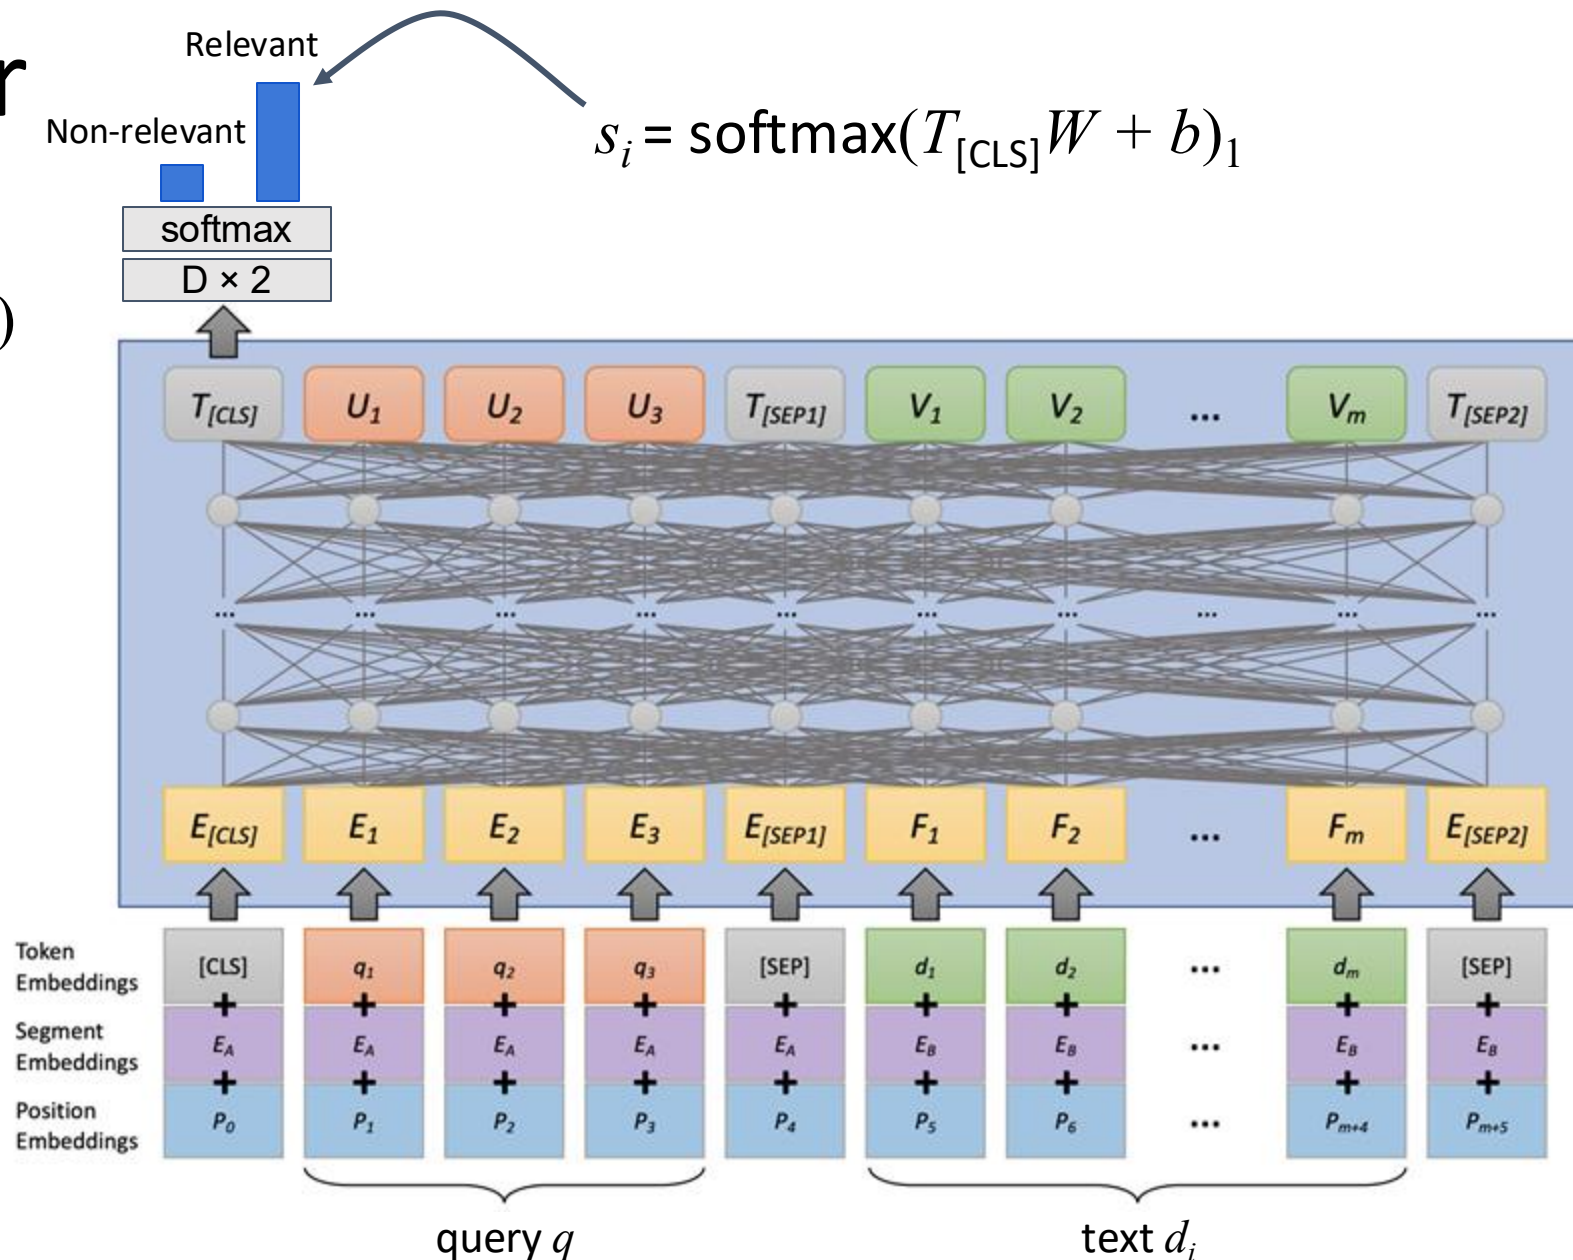
Outline

- Part 1: Background
(text ranking, IR, ML)
- **Part 2: Ranking with relevance classification**
- Part 3: Ranking with dense representations
- Part 4: Conclusion & future directions


monoBERT: BERT reranker

We want:


$$s_i = P(\text{Relevant} = 1 | q, d_i)$$



Once monoBERT is trained...



BERT's Limitations




Cannot input entire documents

- what do we input?
- & how do we label it?

need separate embedding for every possible position
→ restricted to indices 0-511

BERT's Limitations

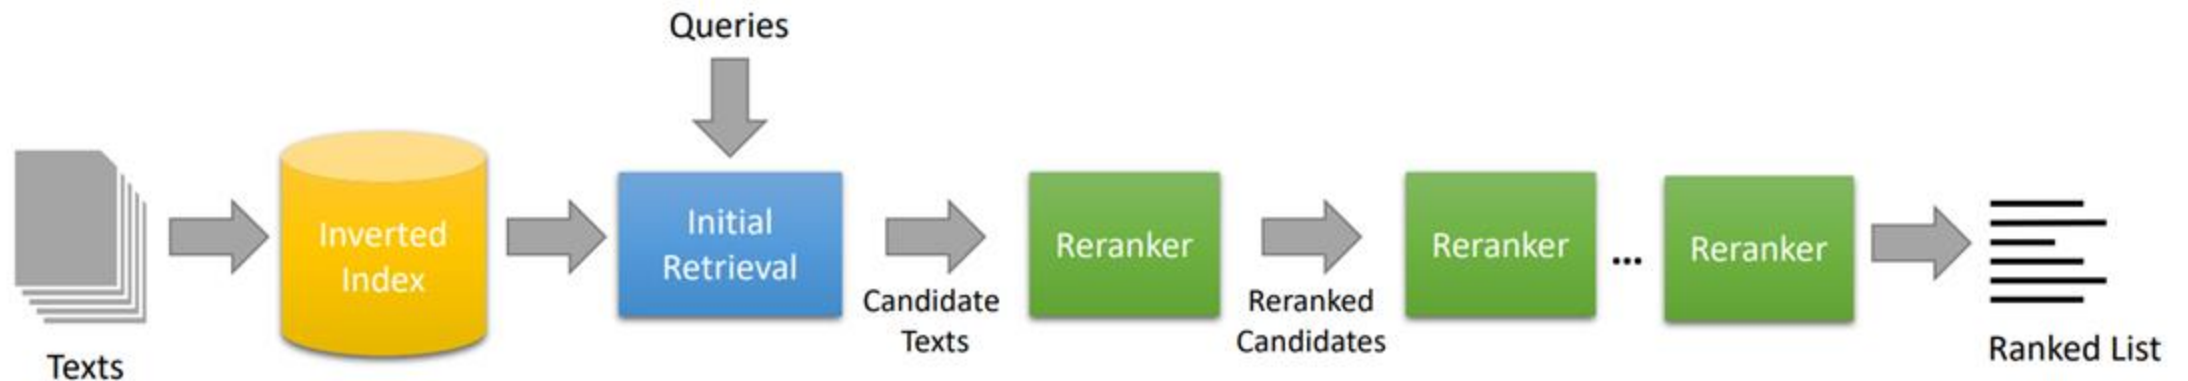
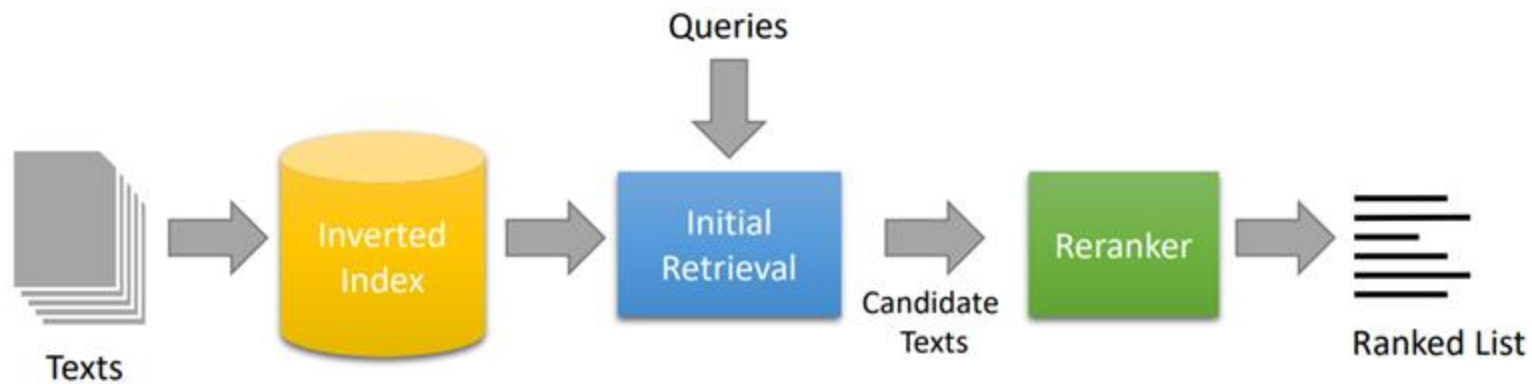


computationally expensive layers
→ e.g., 110+ *million* learned weights
(later: *Beyond BERT & Dense Representations*)

Multi-stage ranking pipeline

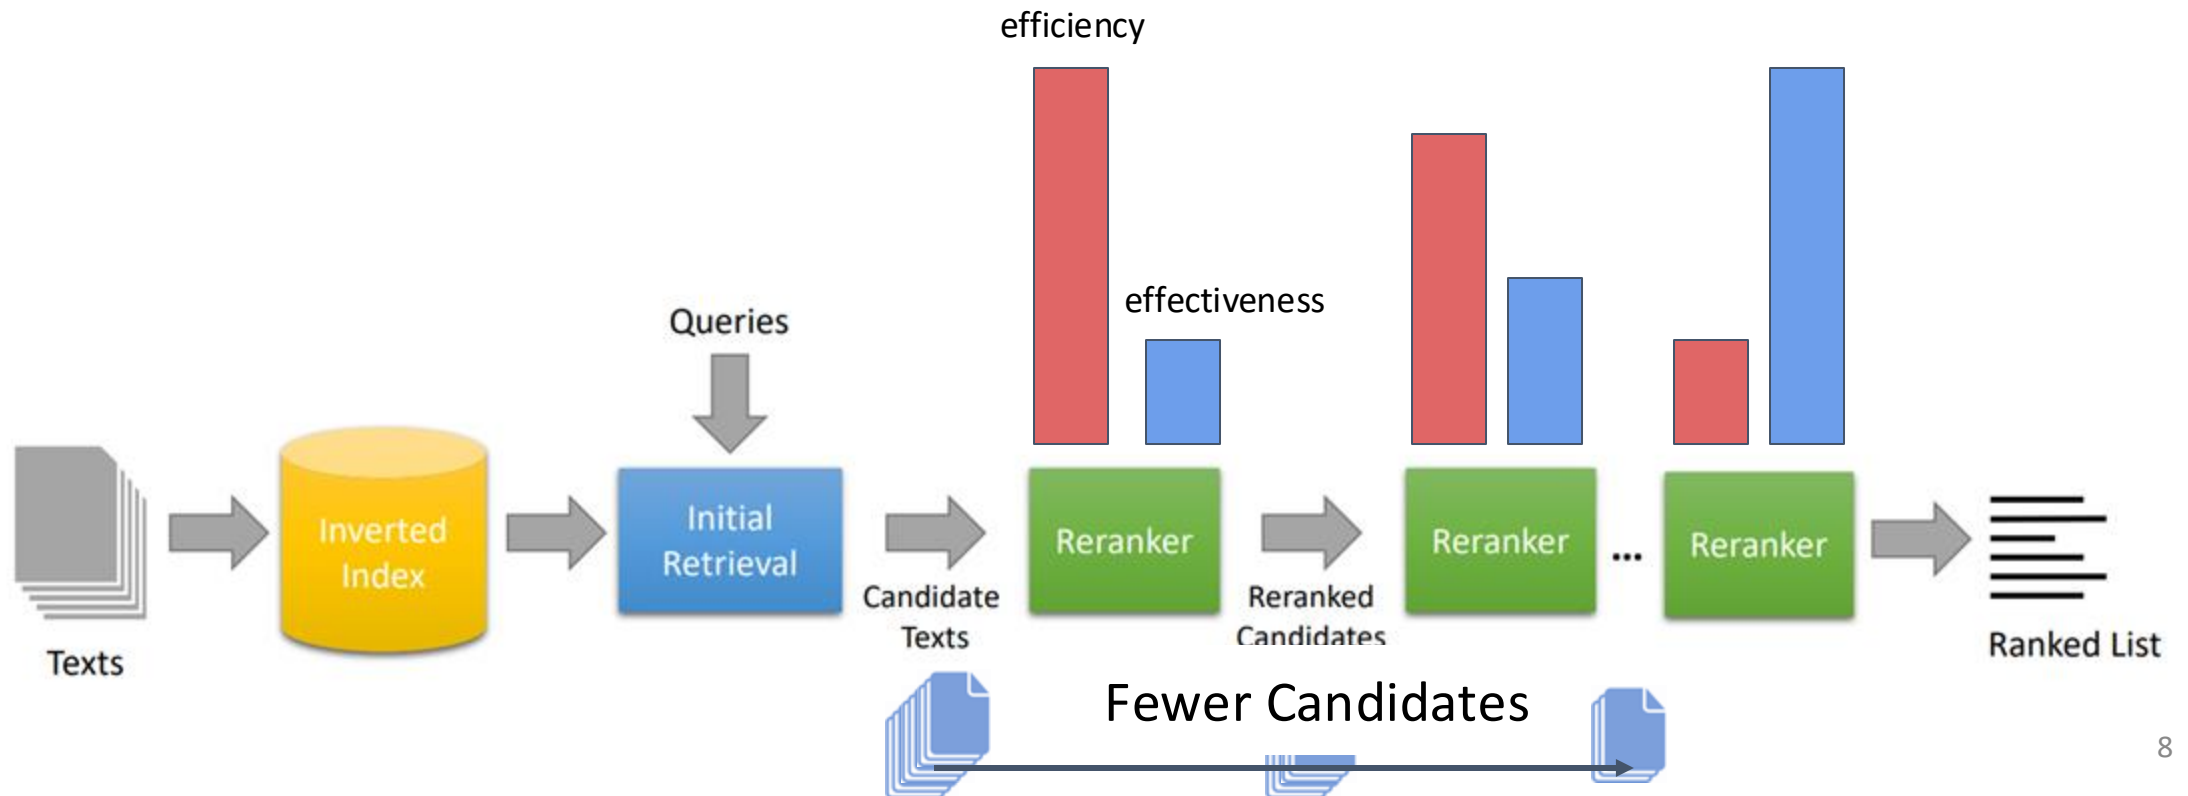
- Identify candidate documents
- Rerank

From Single to Multiple Rerankers

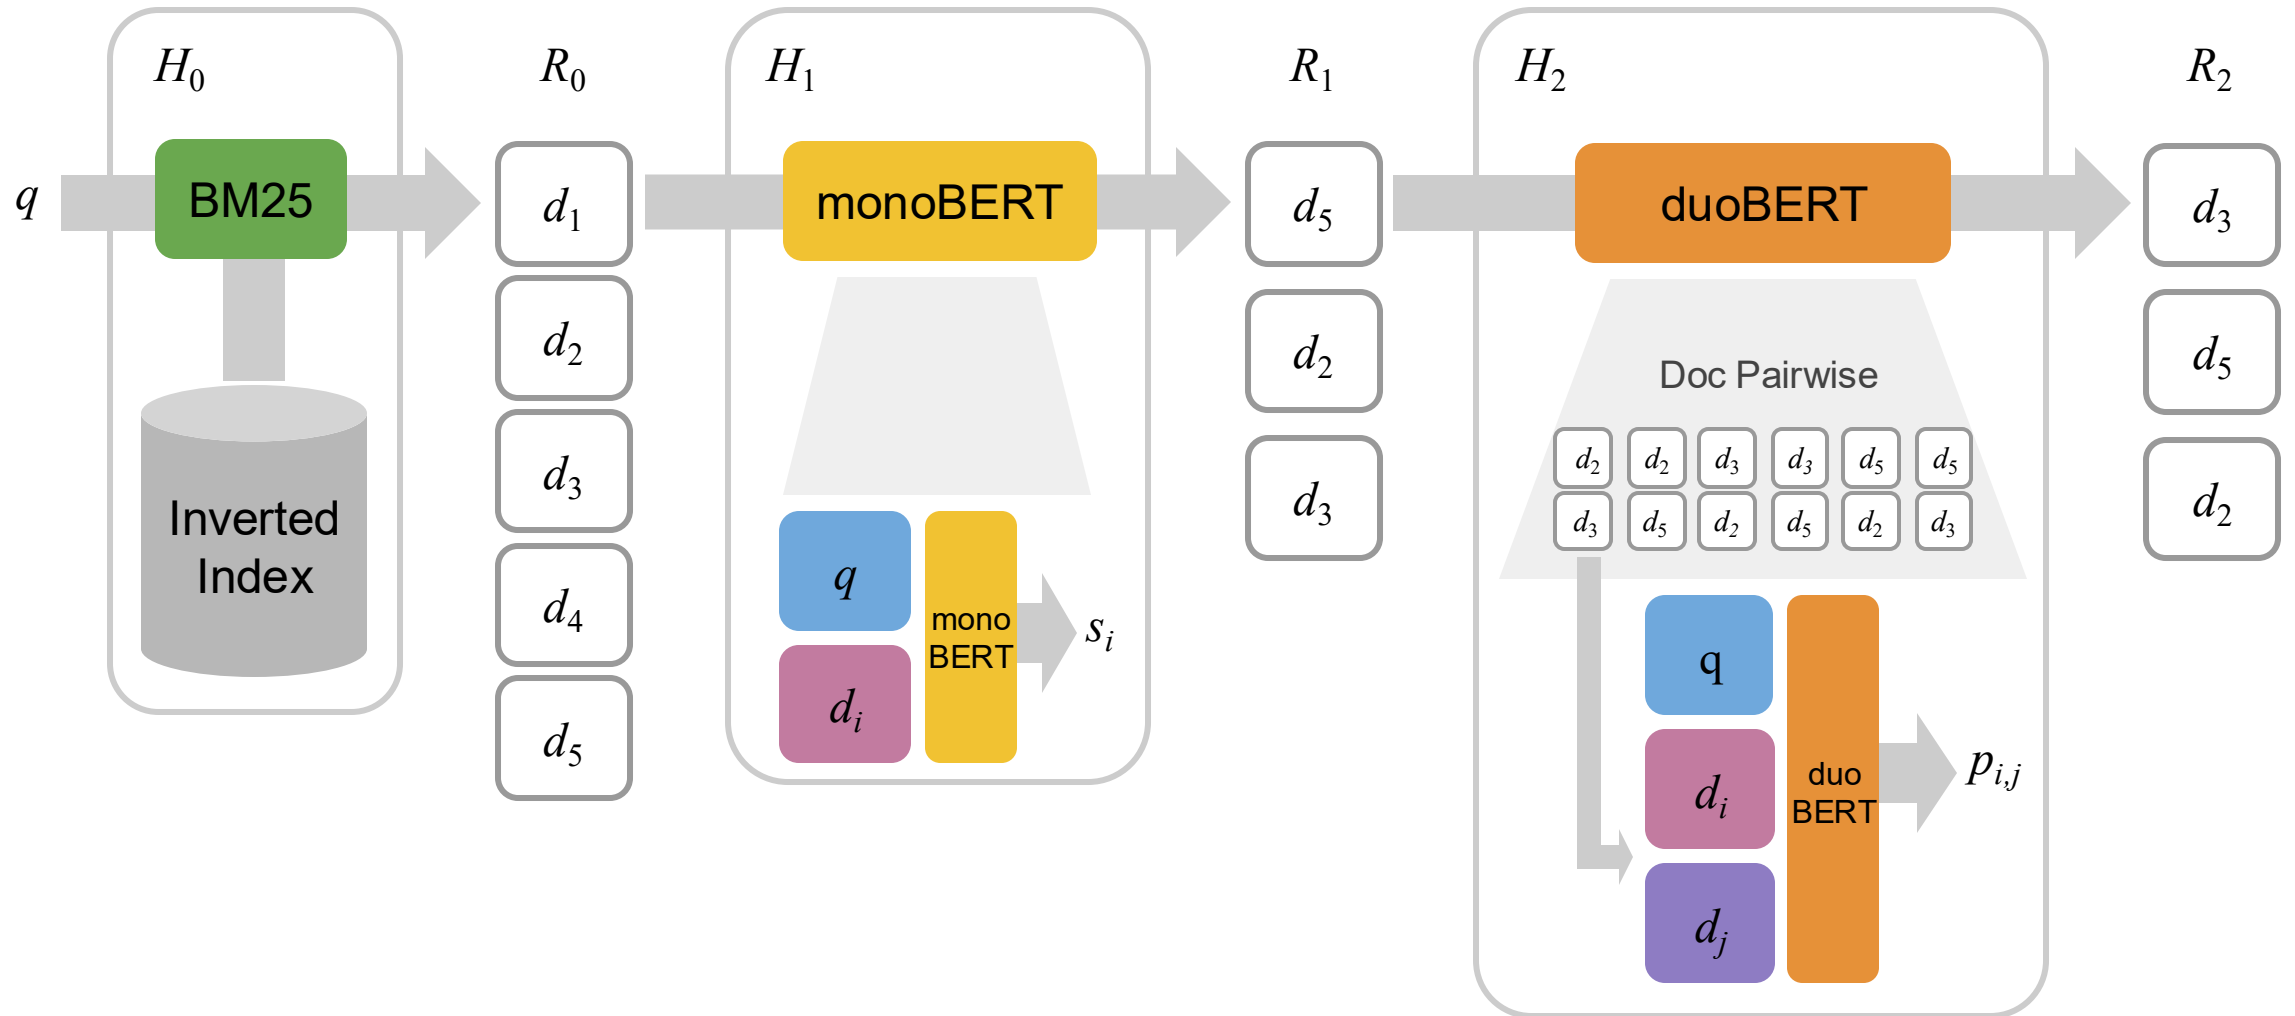


Why Multi-stage?

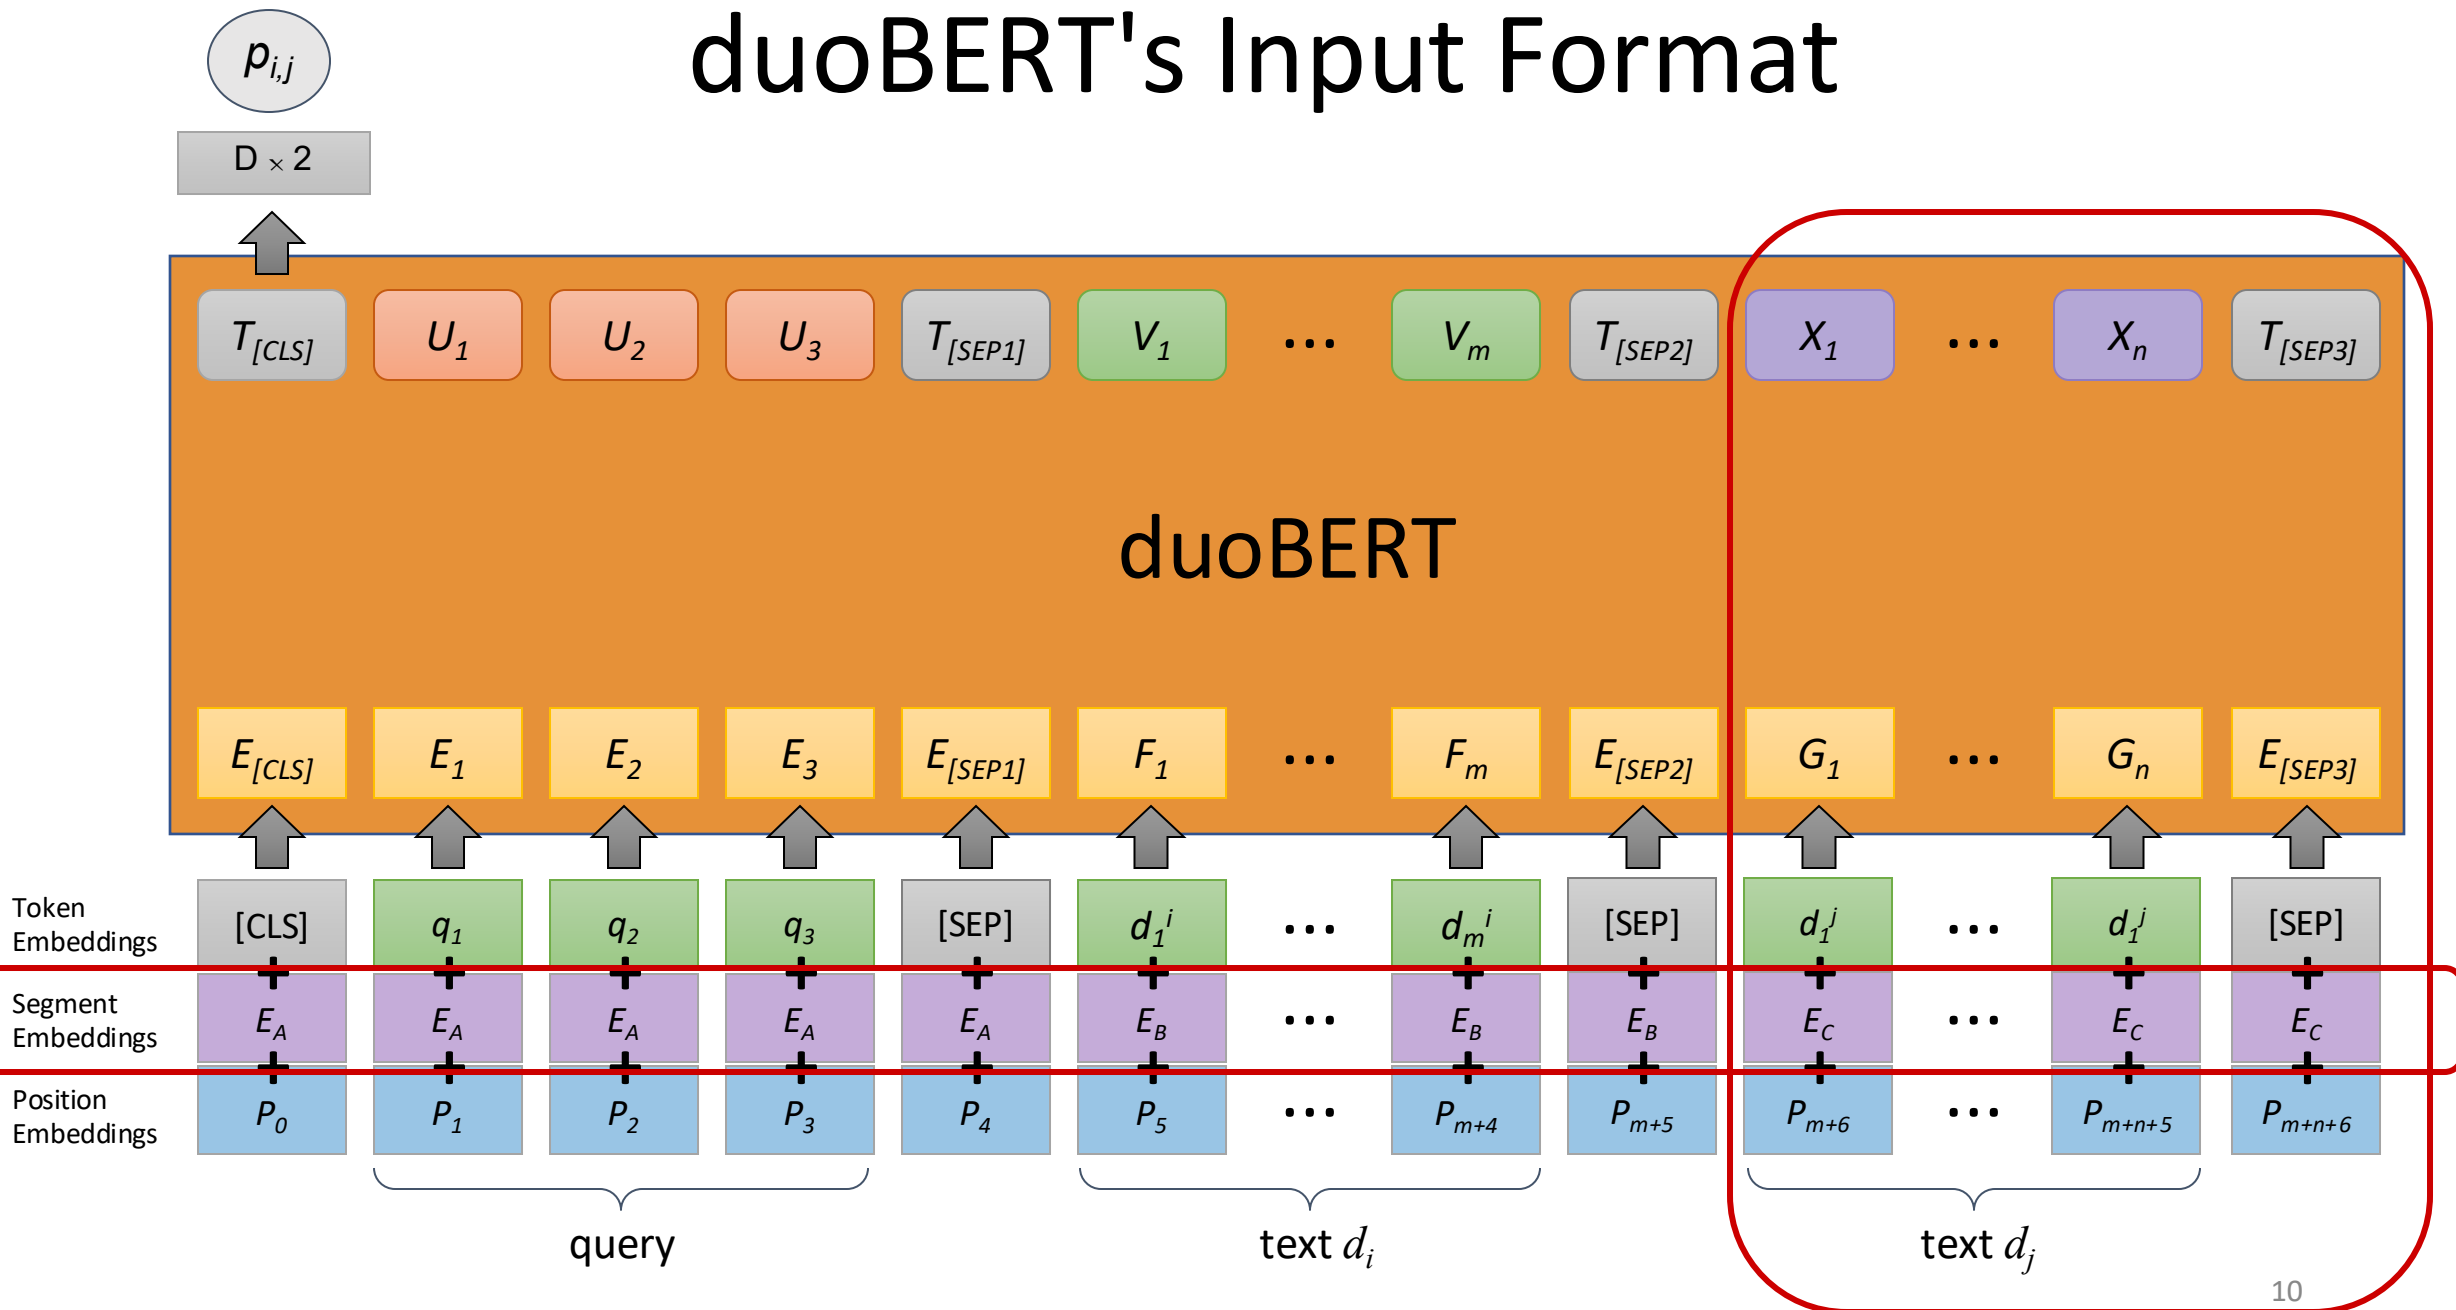
- Trade-off between effectiveness (quality of the ranked lists) and efficiency (retrieval latency)



Multi-stage with duoBERT



duoBERT's Input Format



Training duoBERT

Is doc d_i more relevant than doc d_j to the query q ?

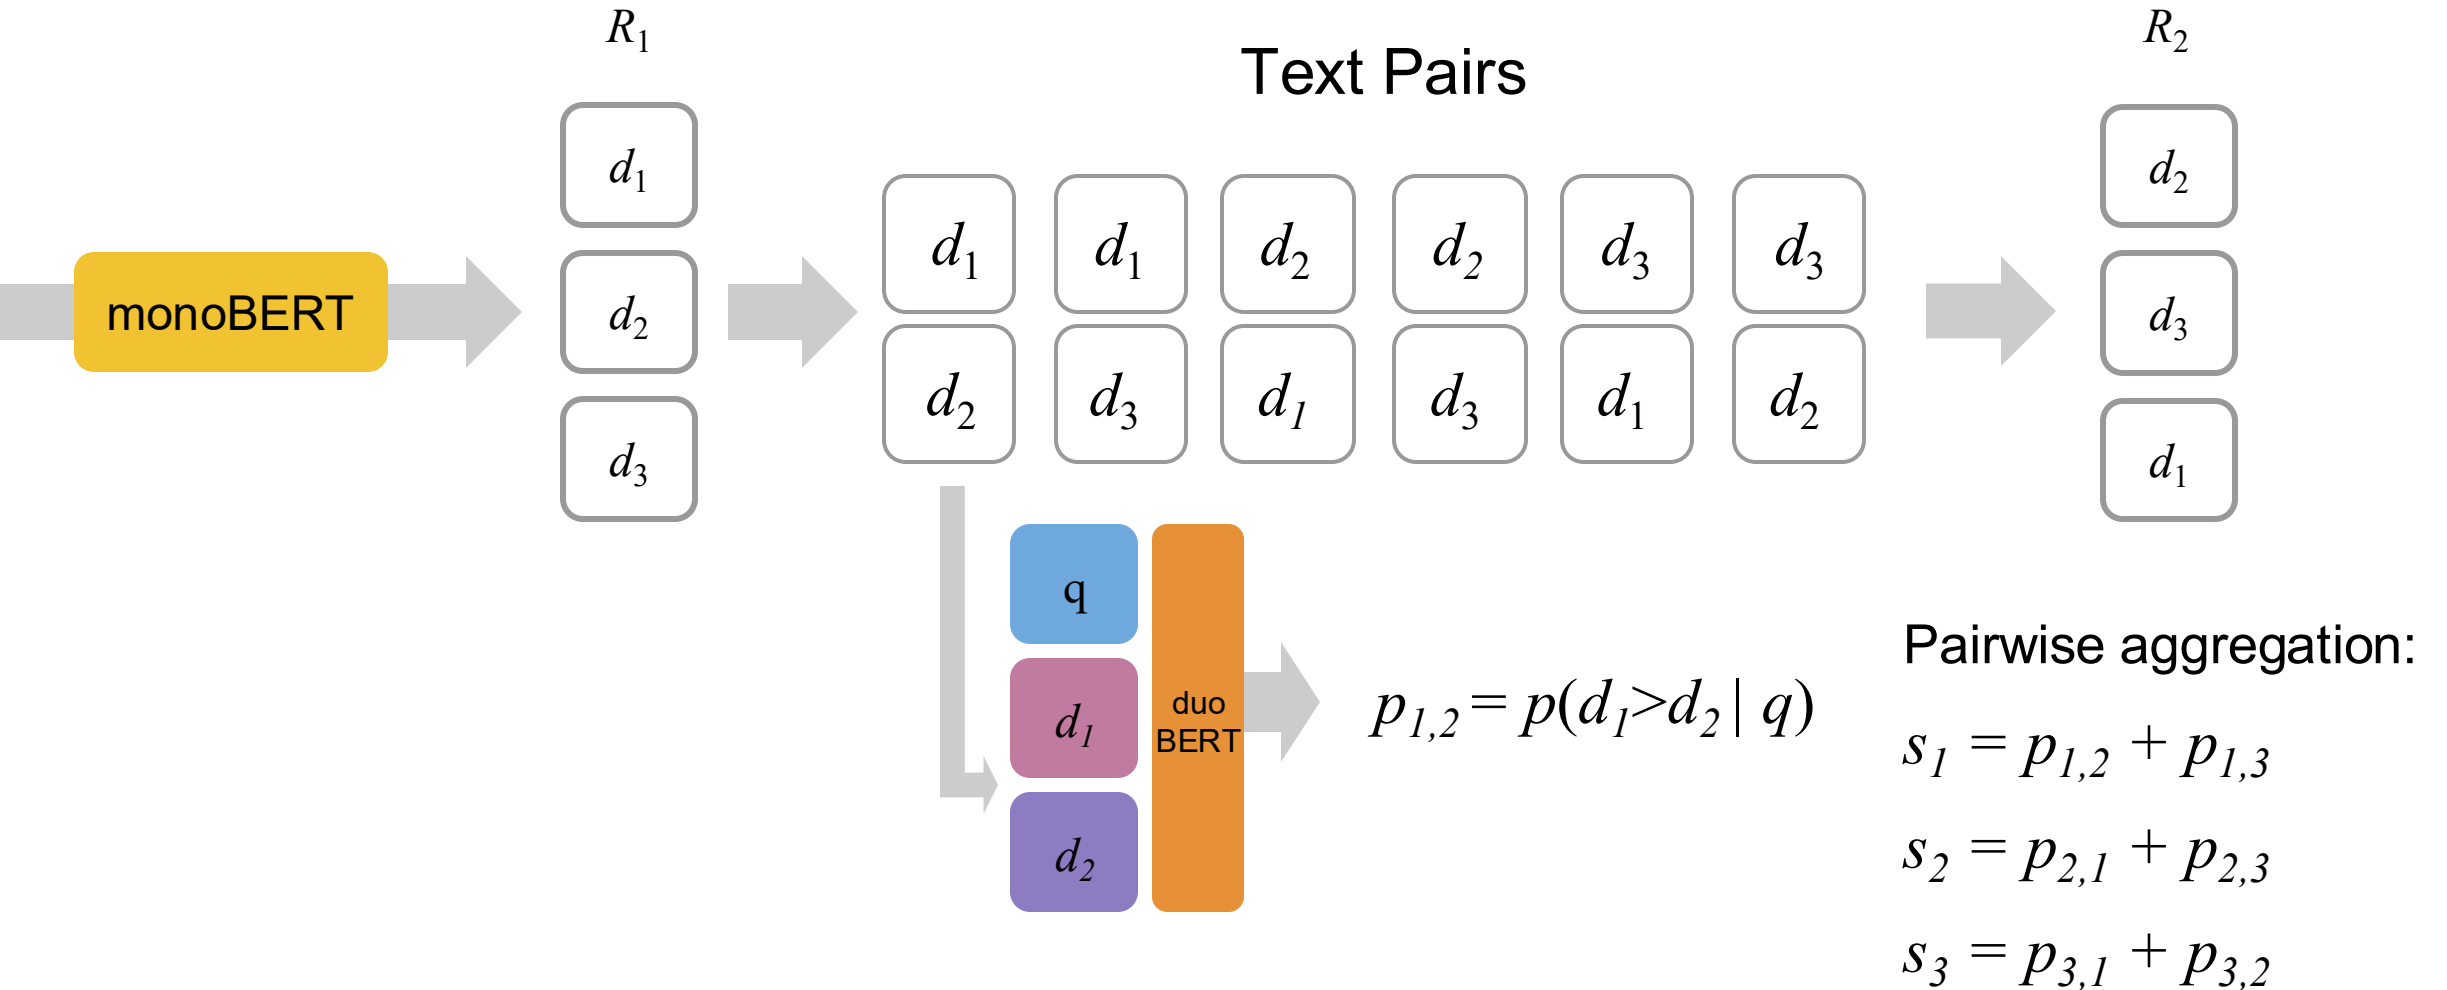
Loss:

$$L_{\text{duo}} = - \sum_{i \in J_{\text{pos}}, j \in J_{\text{neg}}} \log(p_{i,j}) - \sum_{i \in J_{\text{neg}}, j \in J_{\text{pos}}} \log(1 - p_{i,j})$$

$$p_{i,j} = p(d_i > d_j \mid q)$$



Inference with duoBERT



Outline

- Part 1: Background
(text ranking, IR, ML)
- Part 2: Ranking with relevance classification
- **Part 3: Ranking with dense representations**
- Part 4: Conclusion & future directions

Sparse Representations

Task: Estimate the relevance of text d to a query q :

q = "fix my air conditioner"


d = "... AC repair ..."

$$\text{BM25}(q, d) = \sum_{t \in q \cap d} \log \frac{N - \text{df}(t) + 0.5}{\text{df}(t) + 0.5} \cdot \frac{\text{tf}(t, d) \cdot (k_1 + 1)}{\text{tf}(t, d) + k_1 \cdot (1 - b + b \cdot \frac{l_d}{L})}$$

Advantages: 1) Fast to retrieve candidates from a inverted index because q is usually short. 2) Fast to compute because $q \cap d$ is usually small

Disadvantage: Terms need to match exactly

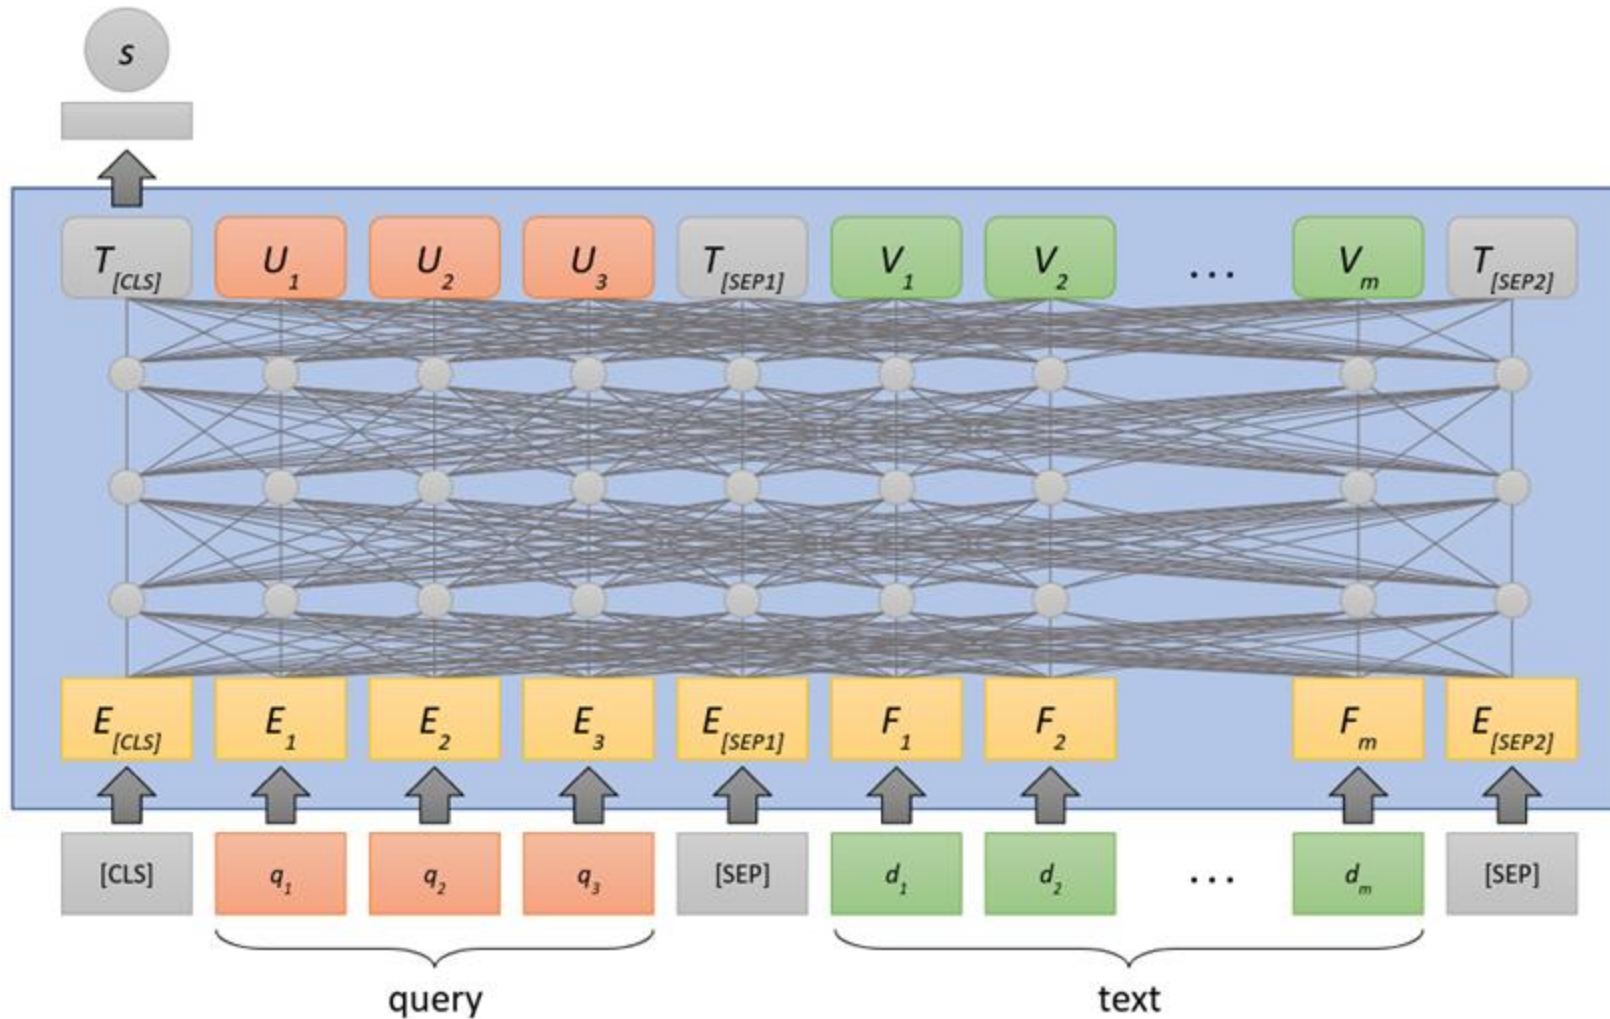
Dense Representations



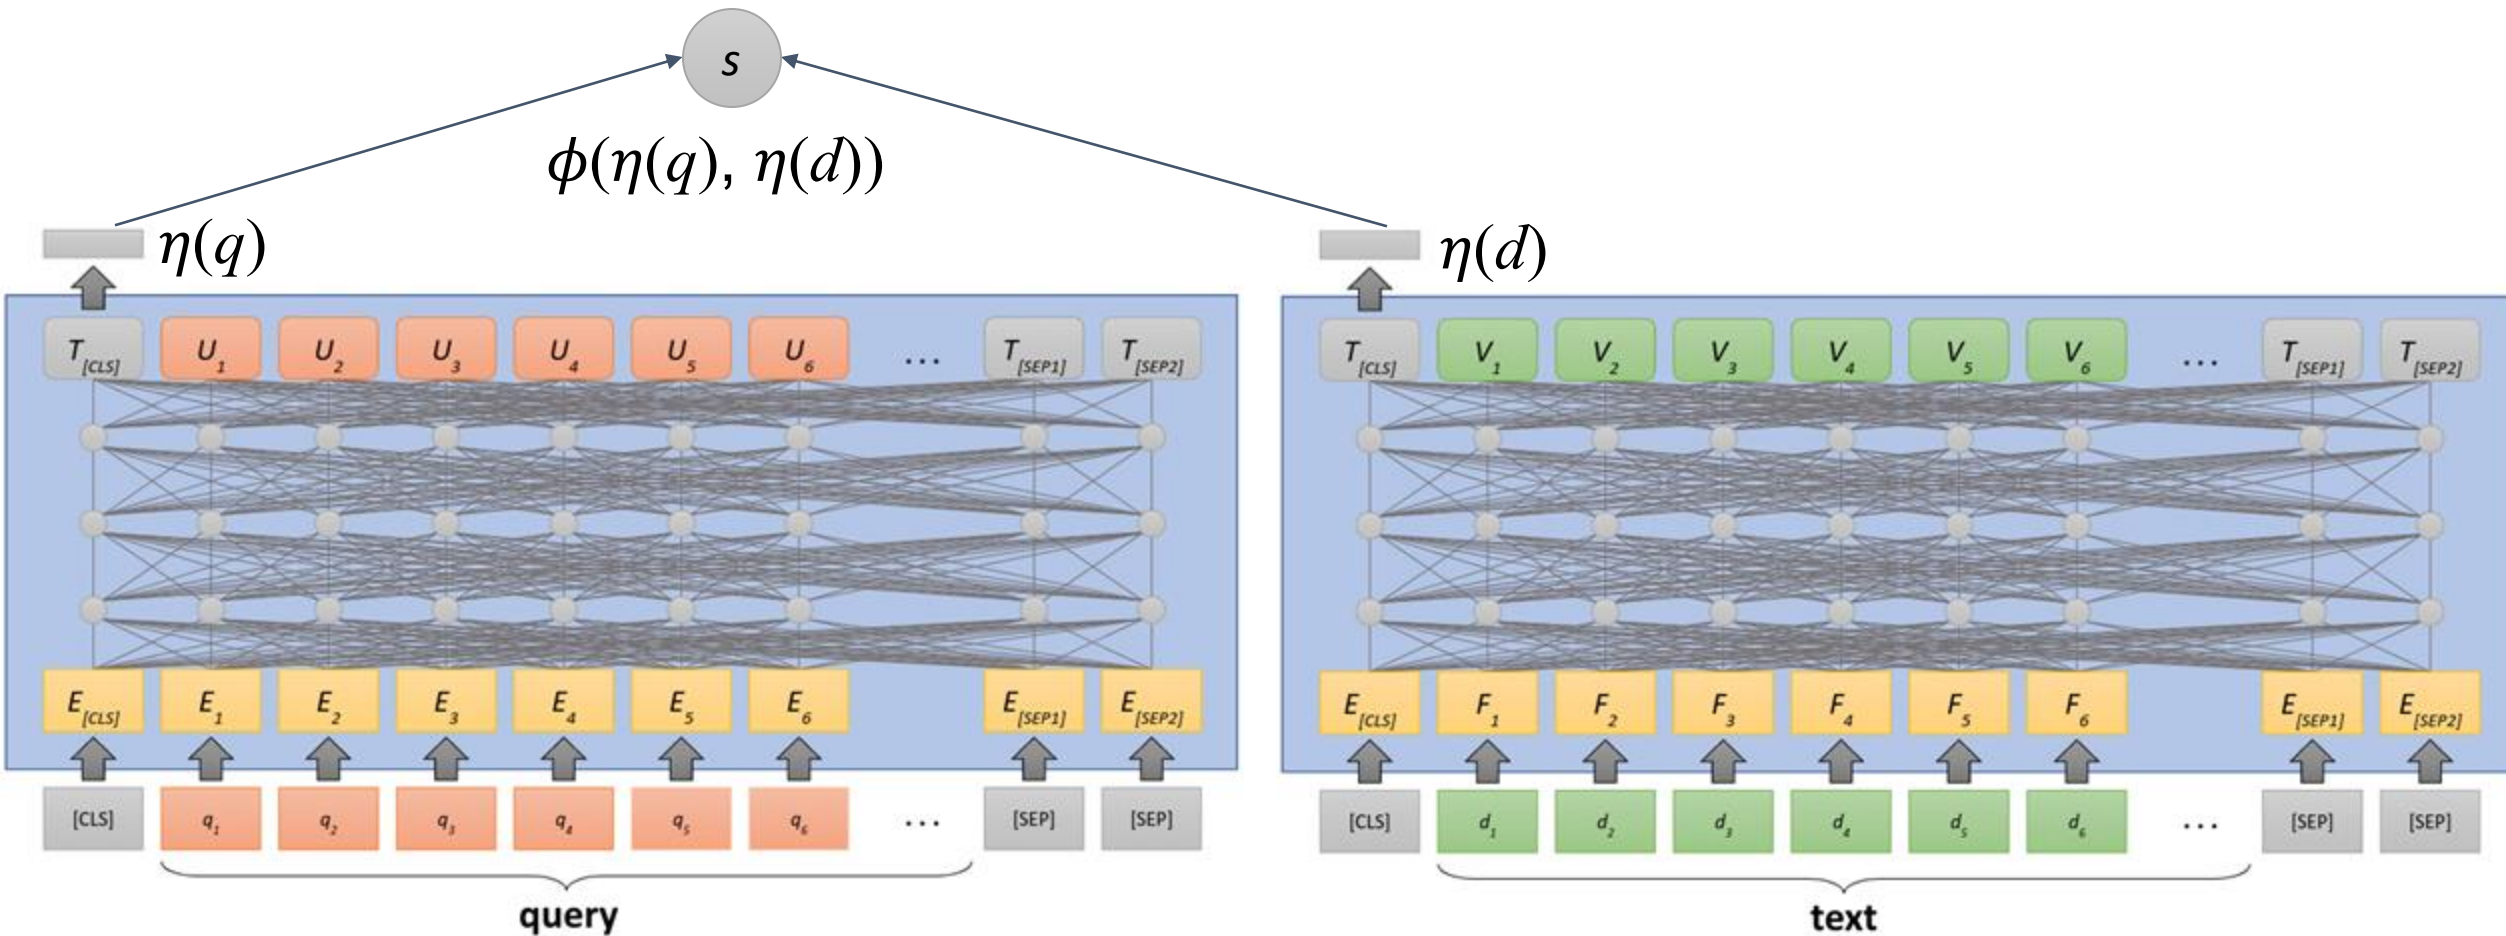
ϕ is a similarity function (e.g., inner product or cosine similarity)

$\phi(\eta(q), \eta(d)) \rightarrow$ ideally measures how relevant q and d are to each other

Types of Encoders: Cross-encoder



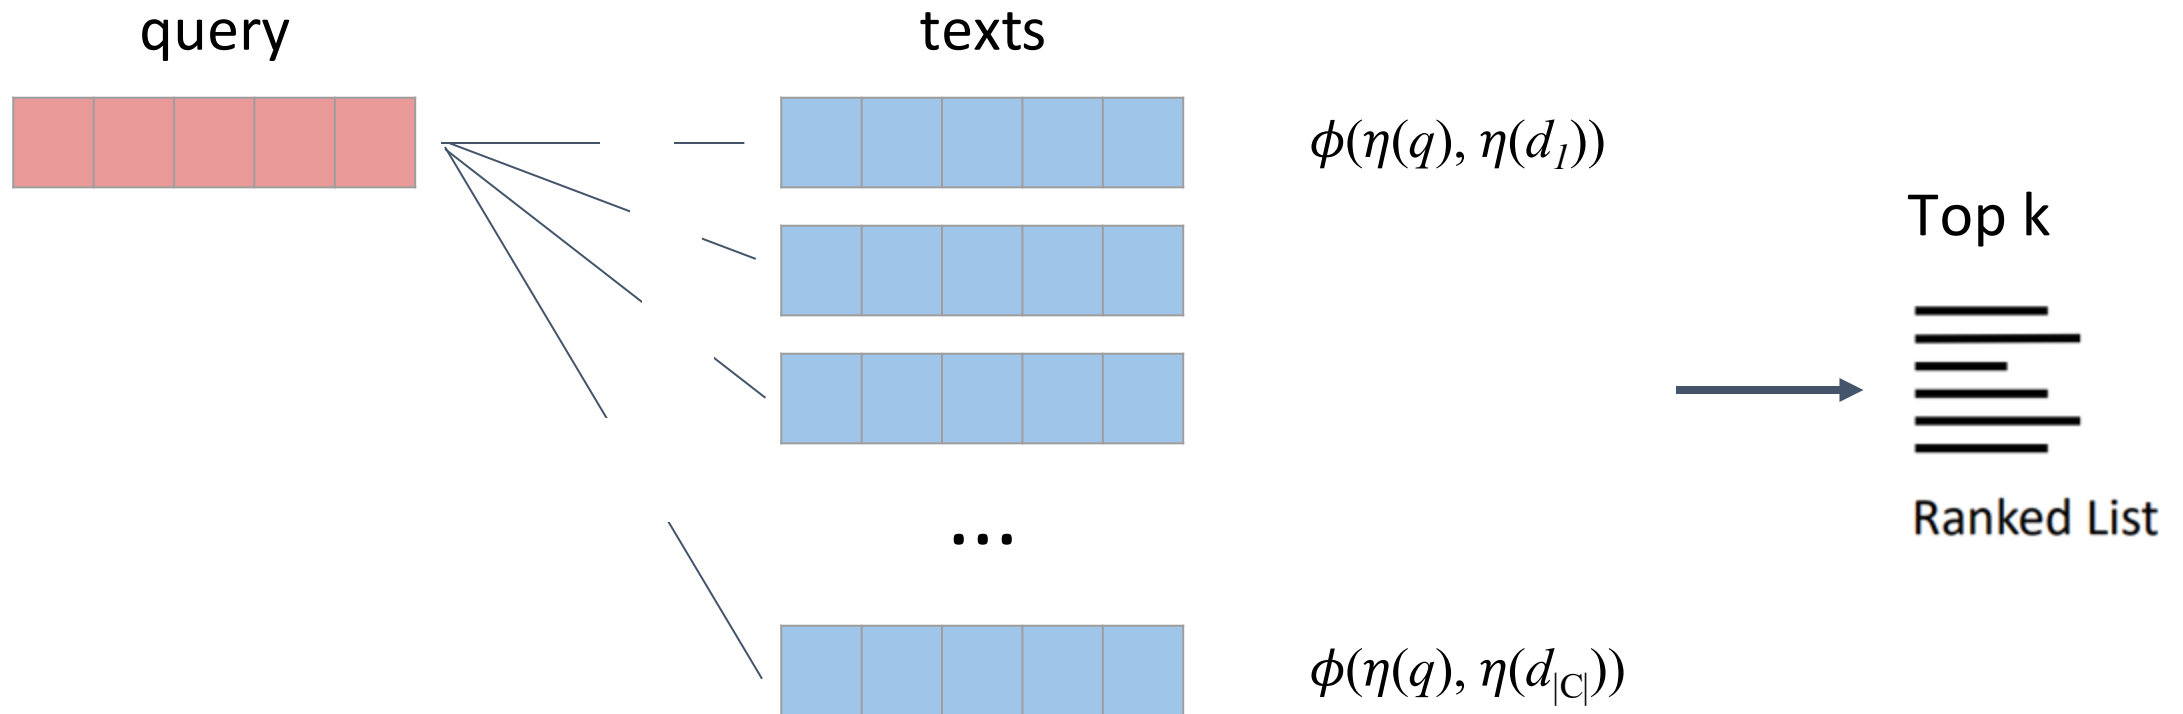
Types of Encoders: Bi-encoder



Nearest Neighbor Search

Task: find the top k most relevant texts to a query

Brute-force search:




We often need to search many (e.g.: billions) of texts

- Brute-force won't scale

Approximate Nearest Neighbor Search

- Exchange accuracy for speed
- E.g.: k-means:



- In practice, ANN implementations are more complicated
- We assume a fast dense retrieval library is available (e.g.: Faiss, Annoy, ScaNN)

Distance-based Transformer Representations

Distance-based Representations

Key characteristic

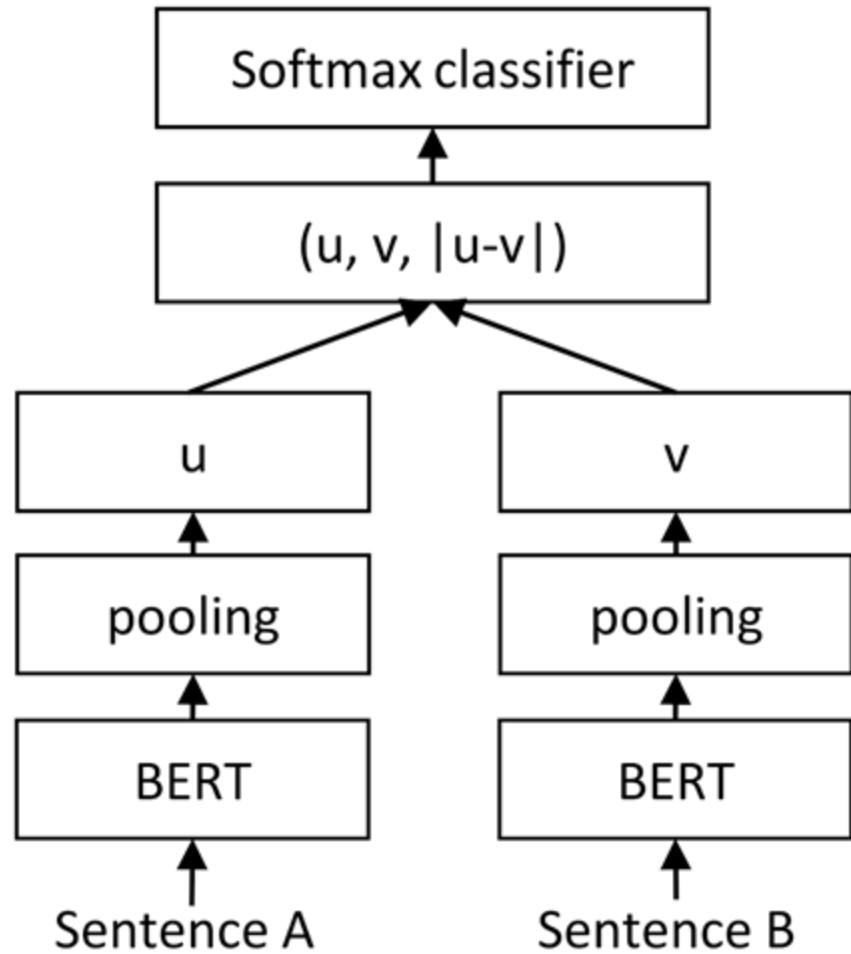
Simple similarity function → inner (dot) product, cosine similarity, ...

$$\phi(u, v) = \eta(u) \cdot \eta(v)$$

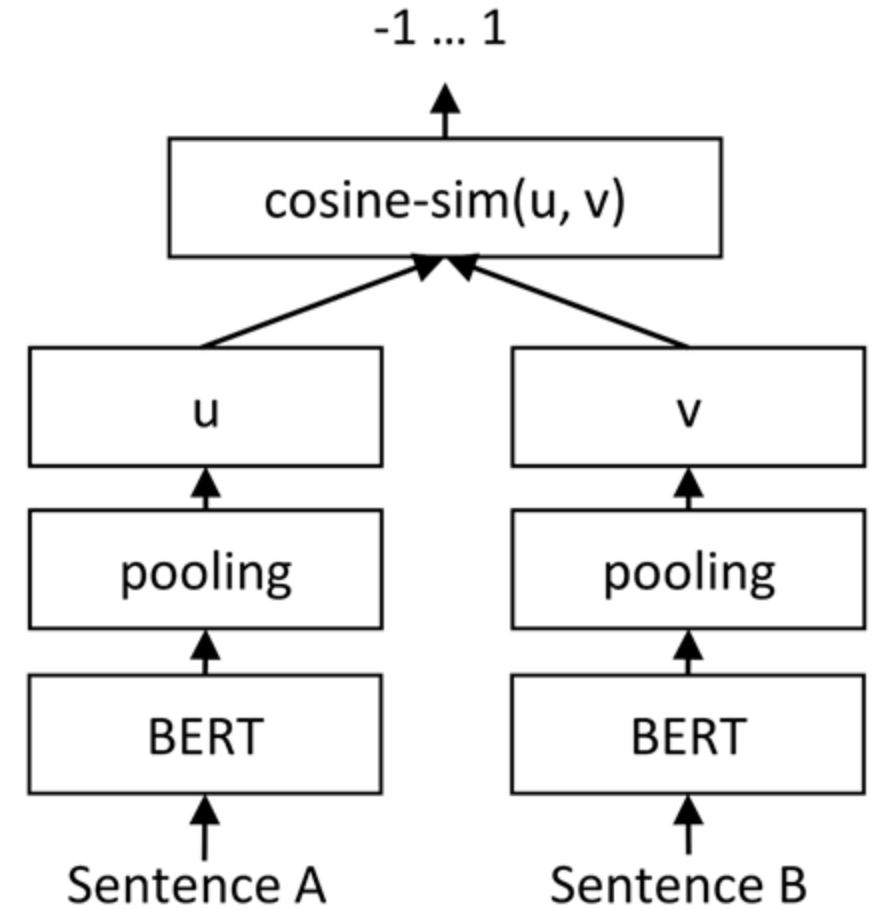
Compatible with ANN search

Johnson, Douze, Jégou. Billion-scale similarity search with GPUs. arXiv 2017.

Distance-based: SentenceBERT



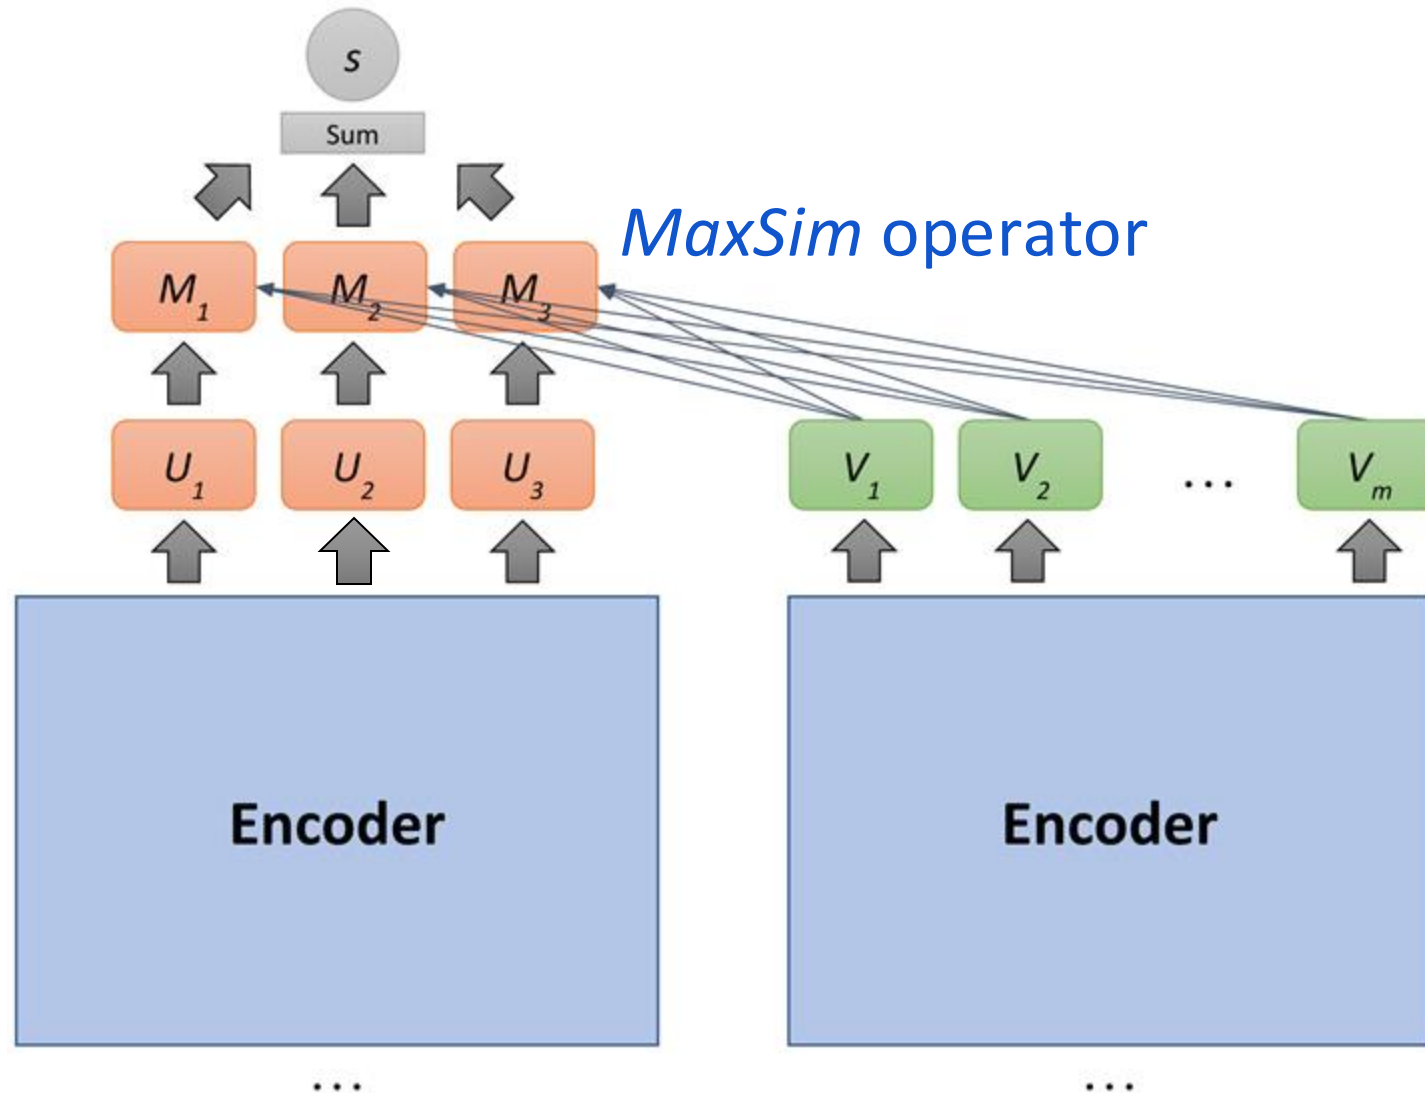
Classification




Regression

Comparison-based Transformer Representations

Comparison-based: ColBERT



Comparison-based: ColBERT




MaxSim:
Sim-mat max pooling
(along query dimension)

$$s_{q,d} = \sum_{i \in \eta(q)} \max_{j \in \eta(d)} \eta(q)_i \cdot \eta(d)_j^T$$

Comparison-based: ColBERT

Compatible with ANN?

- Unclear
- Data-dependent
- 70x faster than BERT-large



Comparison-based: ColBERT

		MS MARCO Passage		
		Development MRR@10	Recall@1k	Latency (ms)
(1a)	BM25 (Anserini, top 1000)	0.187	0.861	62
(1b)	+ monoBERT _{Large}	0.374	0.861	32,900
(2)	FastText + ConvKNRM	0.290	-	90
(3)	doc2query-T5	0.277	0.947	87
(4)	ColBERT (over BERT _{Base})	0.360	0.968	458

Document Preprocessing Techniques


Query vs document expansion

doc2query

DeepCT

DeepImpact

Query reformulation as a translation task




Hard: Input has little information

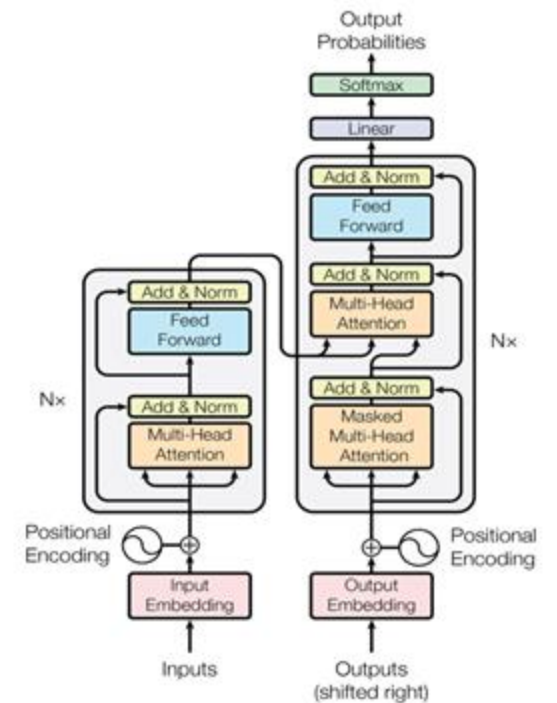


Easier: Input has a lot of information

doc2query

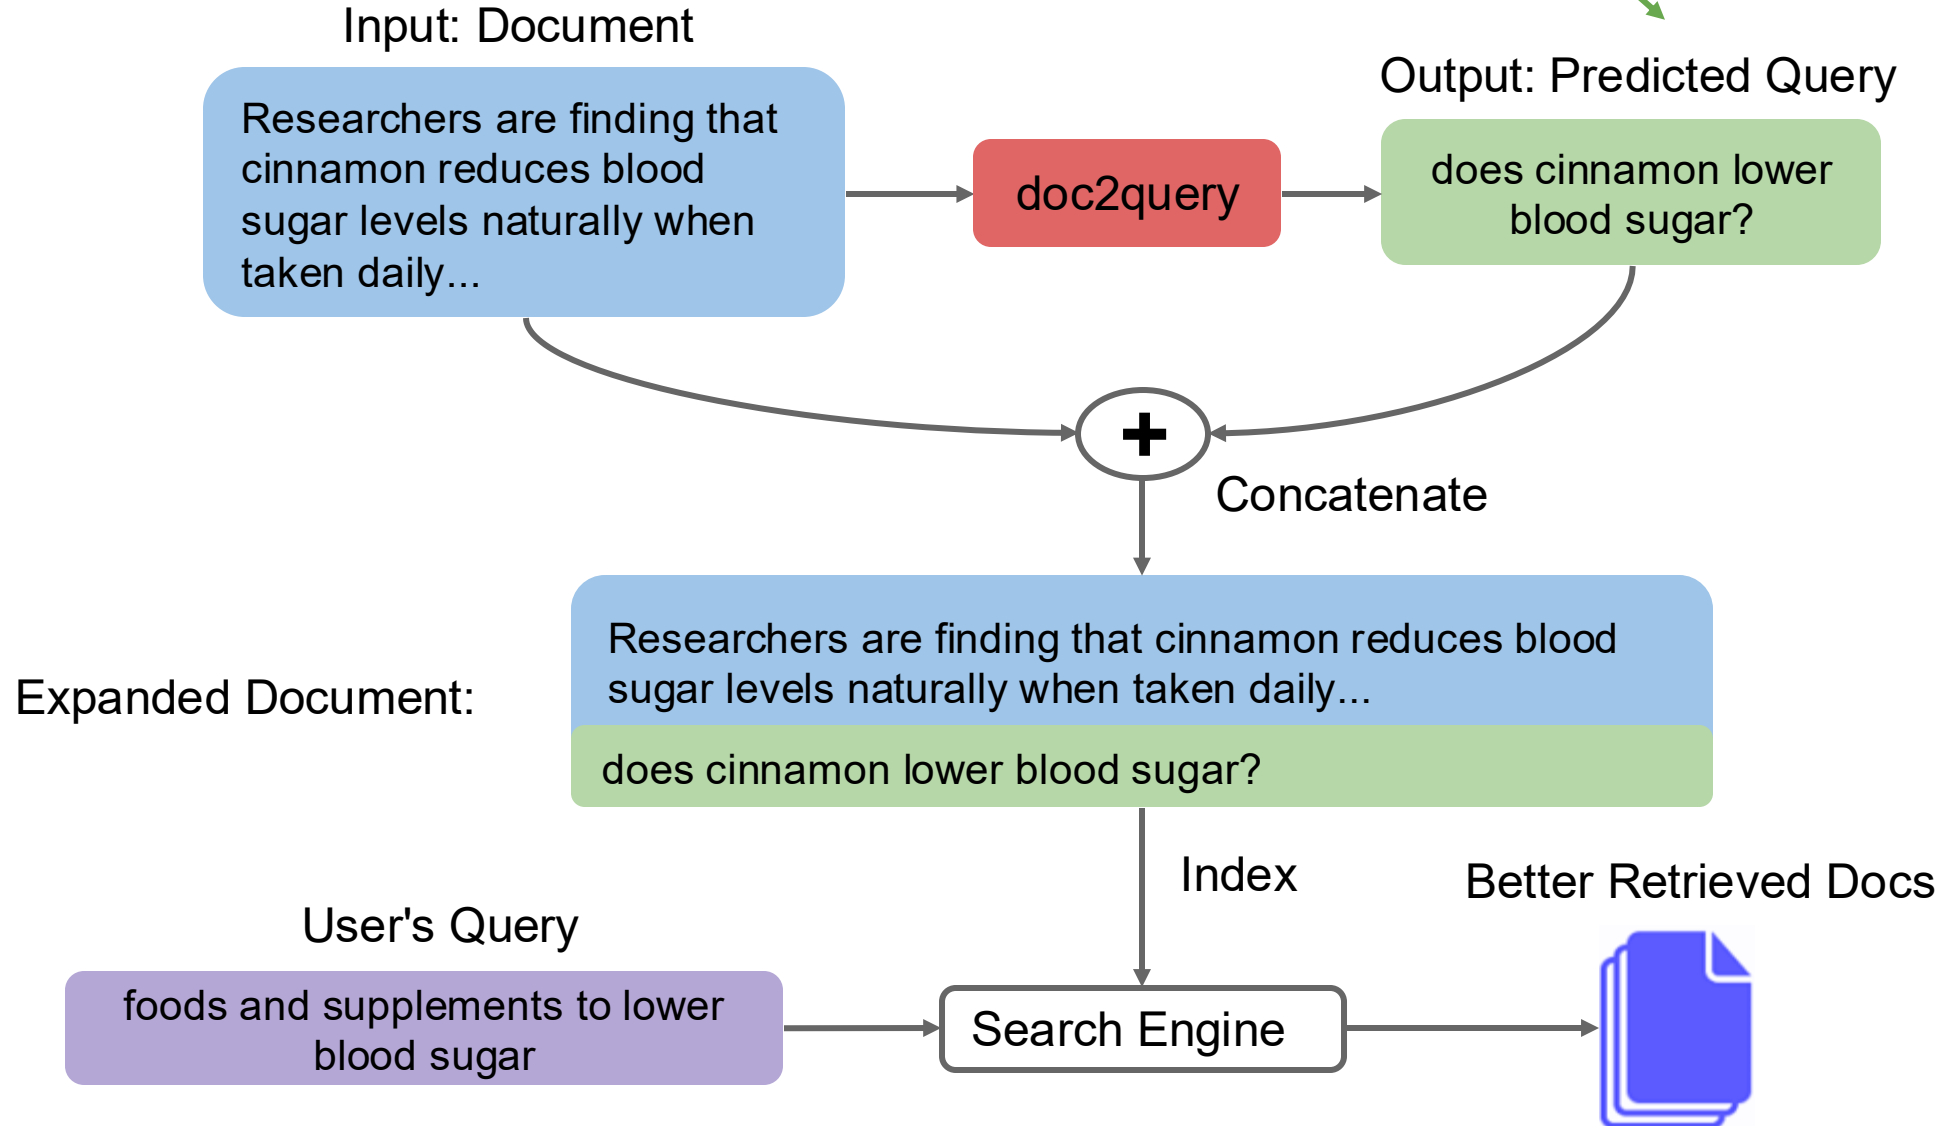


Supervised training:
pairs of <query, relevant document>



doc2query

In practice: 5-40 queries are sampled with top-k or nucleus sampling






Results

	MARCO Passage (MRR@10)	TREC-DL 19 (nDCG@10)	TREC-COVID (nDCG@20)	Robust04 (nDCG@20)
BM25	0.184	0.506	0.659	0.428
+ doc2query	0.277	0.642	0.6375	0.446

zero-shot: doc2query was
trained only on MS MARCO

DeepCT

$$\text{loss} = \sum_t (\hat{y}_{t,d} - y_{t,d})^2$$

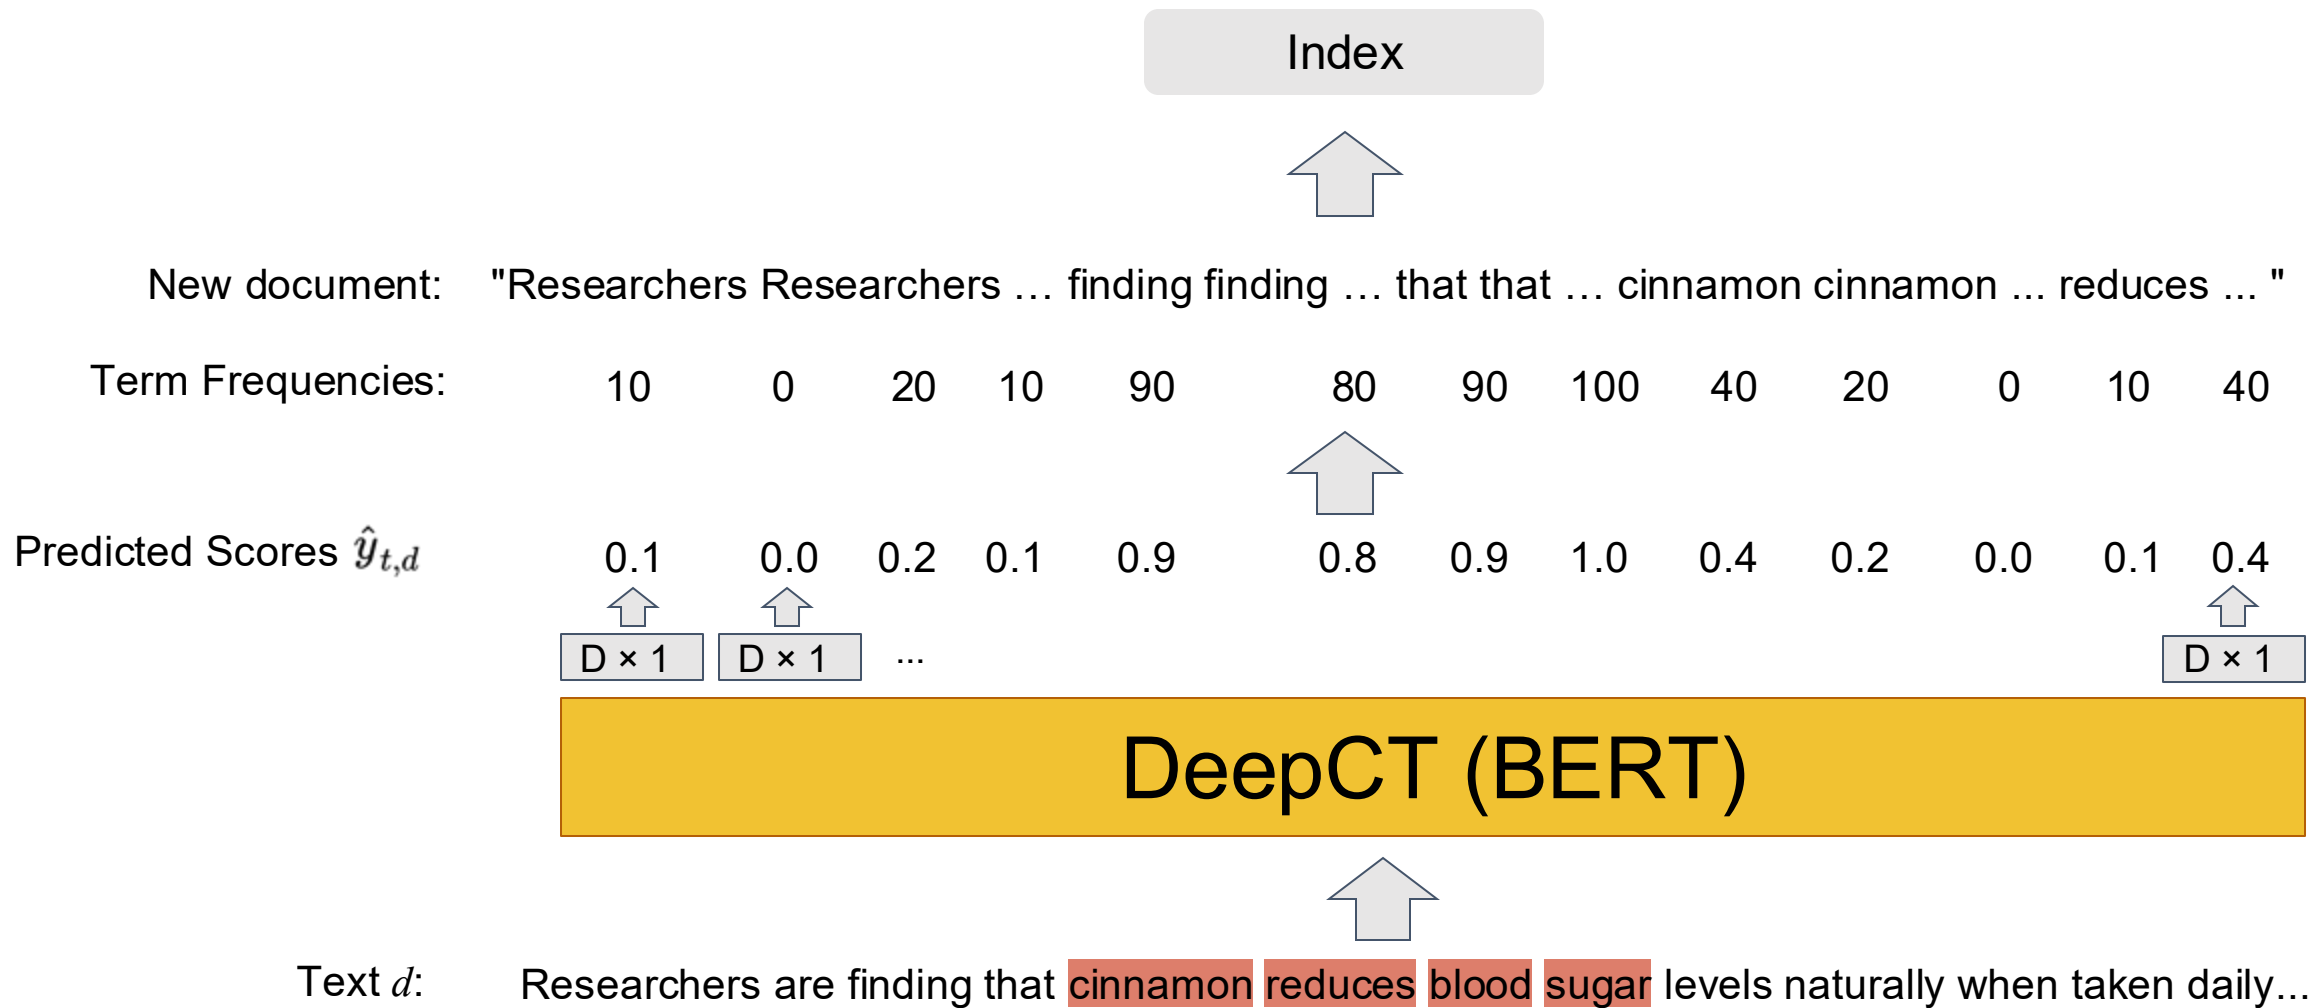
Target Scores	$y_{t,d}$	0.0	1.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
Predicted Scores	$\hat{y}_{t,d}$	0.2	0.5	0.2	0.1	0.4	0.1	0.0	0.6	0.2	0.4	0.3
		 D × 1	 D × 1	...							 D × 1	

DeepCT (BERT)

Text d : The **Geocentric Theory** was proposed by the greeks under the guidance...

Relevant query q : "who proposed the **geocentric theory**"

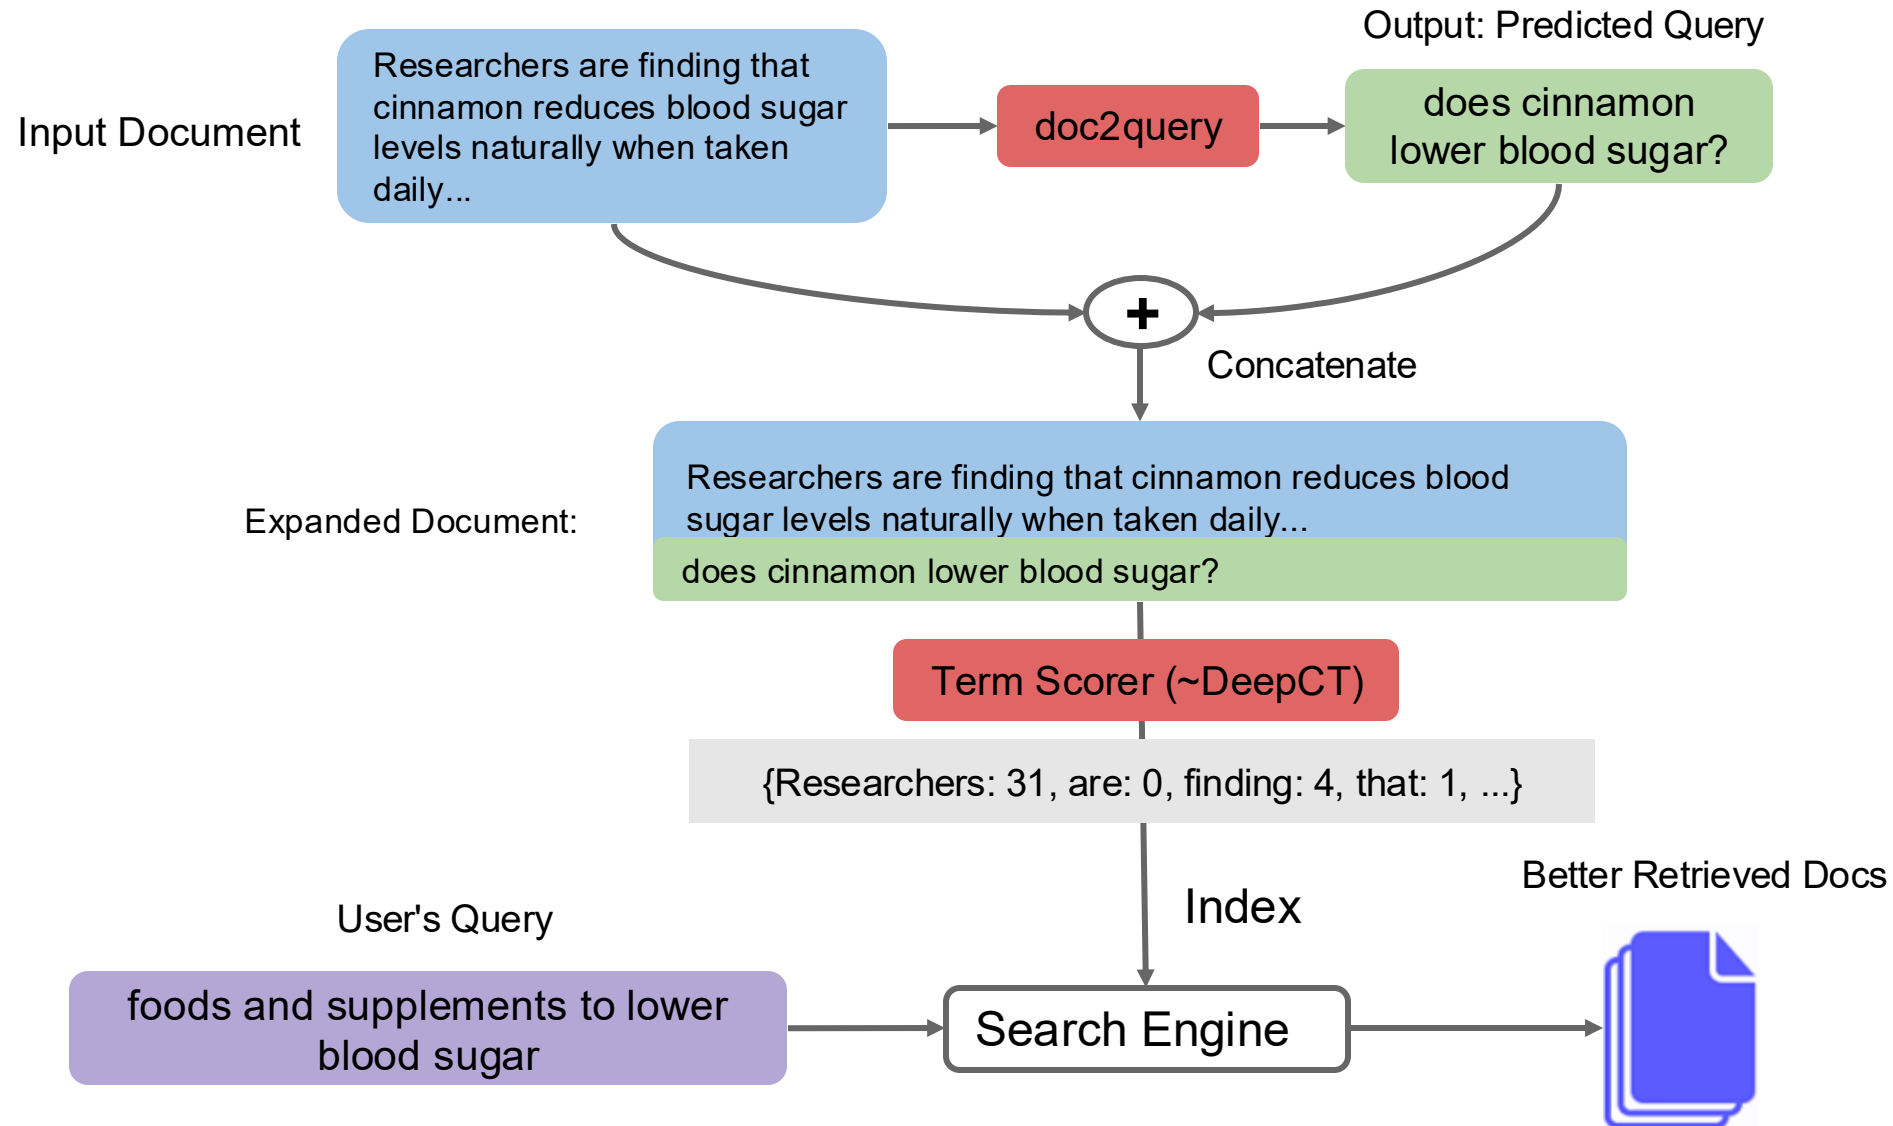
Once DeepCT is trained...



Results on MS MARCO Passage Dev Set

Model	MRR@10	R@1000	BERT Inferences per doc
BM25	0.184	0.853	-
+ doc2query	0.229	0.907	1
+ doc2query	0.277	0.944	40
DeepCT	0.243	0.913	1

DeepImpact: combining doc2query with DeepCT



Results on MS MARCO Passage Dev Set

Model	MRR@10	R@1000	Latency (ms/query)
BM25	0.184	0.853	13
DeepCT	0.243	0.913	11
doc2query	0.278	0.947	12
DeepImpact	0.326	0.948	58
BM25 + monoBERT	0.355	0.853	(GPU) 10,700

Takeaways of Document Expansion

Advantages:

- Documents have more context than queries → easy prediction task
- Documents can be processed offline *and* in parallel
- Run on CPU at query time

Disadvantages:

- Have to iterate over the entire collection
- Not as effective as rerankers (yet)

Conclusions and Future Directions

Conclusions

- Pretrained Transformers showed significant improvements in various IR benchmarks
- Reproduced and adopted by many in academia and industry
- No doubt we are in the age of BERT and Transformers

Learn more in survey (& upcoming book):

*Pretrained Transformers for Text
Ranking:
BERT and Beyond*

by Jimmy Lin, Rodrigo Nogueira, and Andrew Yates

<https://arxiv.org/abs/2010.06467>

Thanks!